



# PRINCIPLES AND TECHNOLOGIES OF DISTRIBUTED DATABASE SYSTEM – INTRODUCTION

**李旭东** **Li-Xudong**

LEEXUDONG@NANKAI.EDU.CN

NANKAI UNIVERSITY

# OBJECTIVES


- The Age of Database
  - Distributed Data Processing
  - What is a Distributed Database System?
  - Data Delivery Alternatives
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD

# OBJECTIVES

- The Age of Database
  - Distributed Data Processing
  - What is a Distributed Database System?
  - Data Delivery Alternatives
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD

# THE AGE OF DATABASE

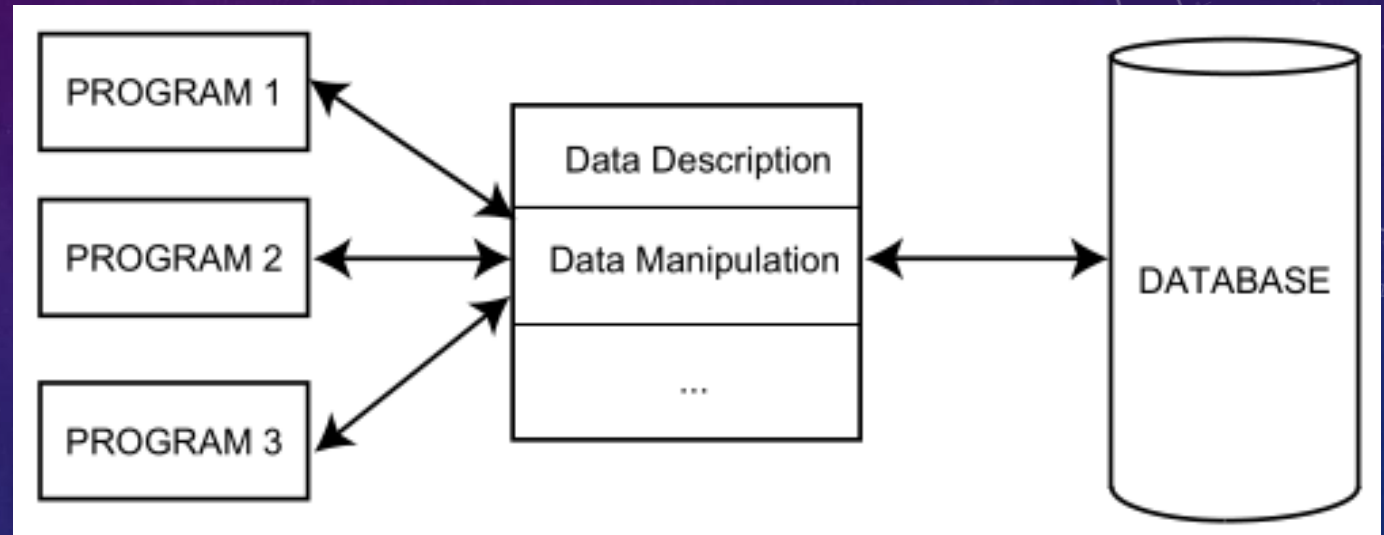
- Data
- File system





# THE AGE OF DATABASE

- Data
- File system
- Network
- Data Independence(数据独立性)
- Database system



# DATABASE SYSTEM

- Database-management system
  - DBMS
  - A collection of interrelated data and a set of programs to access those data
- Goal of DBMS
  - Provide a way to store and retrieve database information that is both convenient and efficient

# DATABASE-SYSTEM APPLICATIONS

- Enterprise Information
  - Sales, accounting, human resources, manufacturing, online retailers
- Banking and Finance
  - Banking, credit card transactions, finance stock
- Universities
- Airlines
- Telecommunication

# PURPOSE OF DATABASE SYSTEMS

- Enable durability (持久性) storage
- Large amounts of data
- Isolation (独立性) among users
- Atomicity(原子性)
- Reduce redundancy(冗余)
- Consistency(一致性)
- Integrity (完整性)
- Efficient store and access
- Concurrent(并发) access
- Security(安全)



# DATA : CASE

```
type instructor = record
```

```
  ID : string;
```

```
  name : string;
```

```
  dept_name : string;
```

```
  salary : integer;
```

```
  birth : date;
```

```
  gender : boolean;
```

```
end;
```

©LXD

```
type department = record
```

```
  ...
```

```
end;
```

```
type student = record
```

```
  ...
```

```
end;
```

# DATA MODEL

- A collection of tools for describing
  - Data
  - Data relationships
  - Data semantics
  - Data constraints

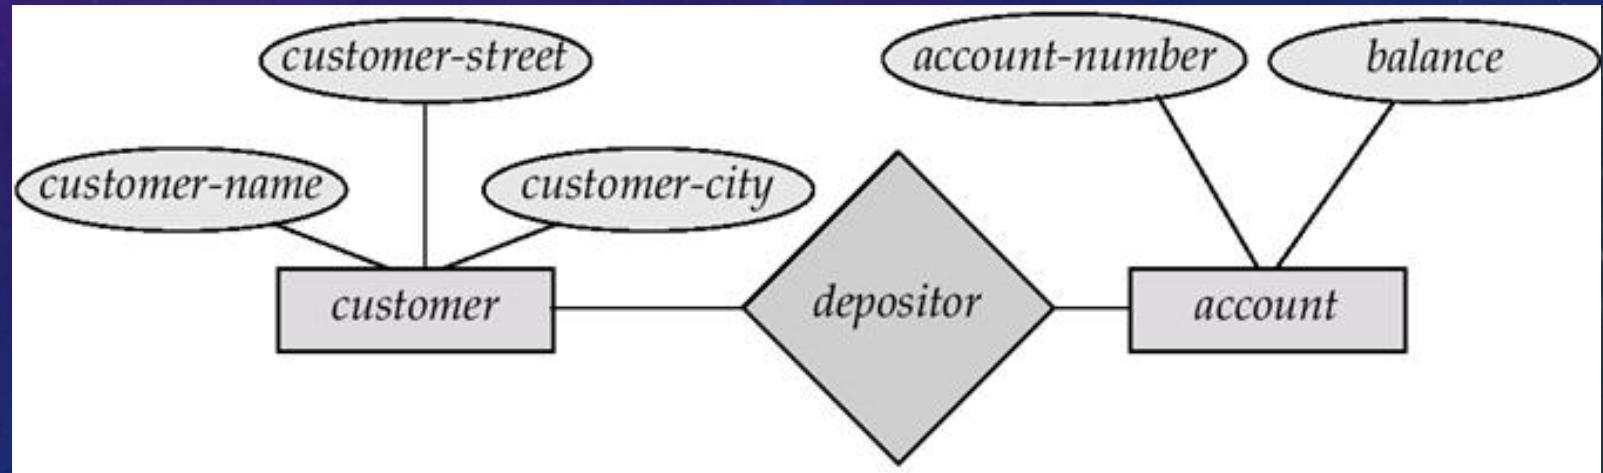
# DATA (DATABASE) MODEL

- Entity-relationship model(实体关系E-R)
- Hierarchical data model(层次)
- Network data model(网络)
- Relational data model(关系)
- Object-based data model(对象)
- Semi-structured data model (半结构)
- Graph data model(图), ...



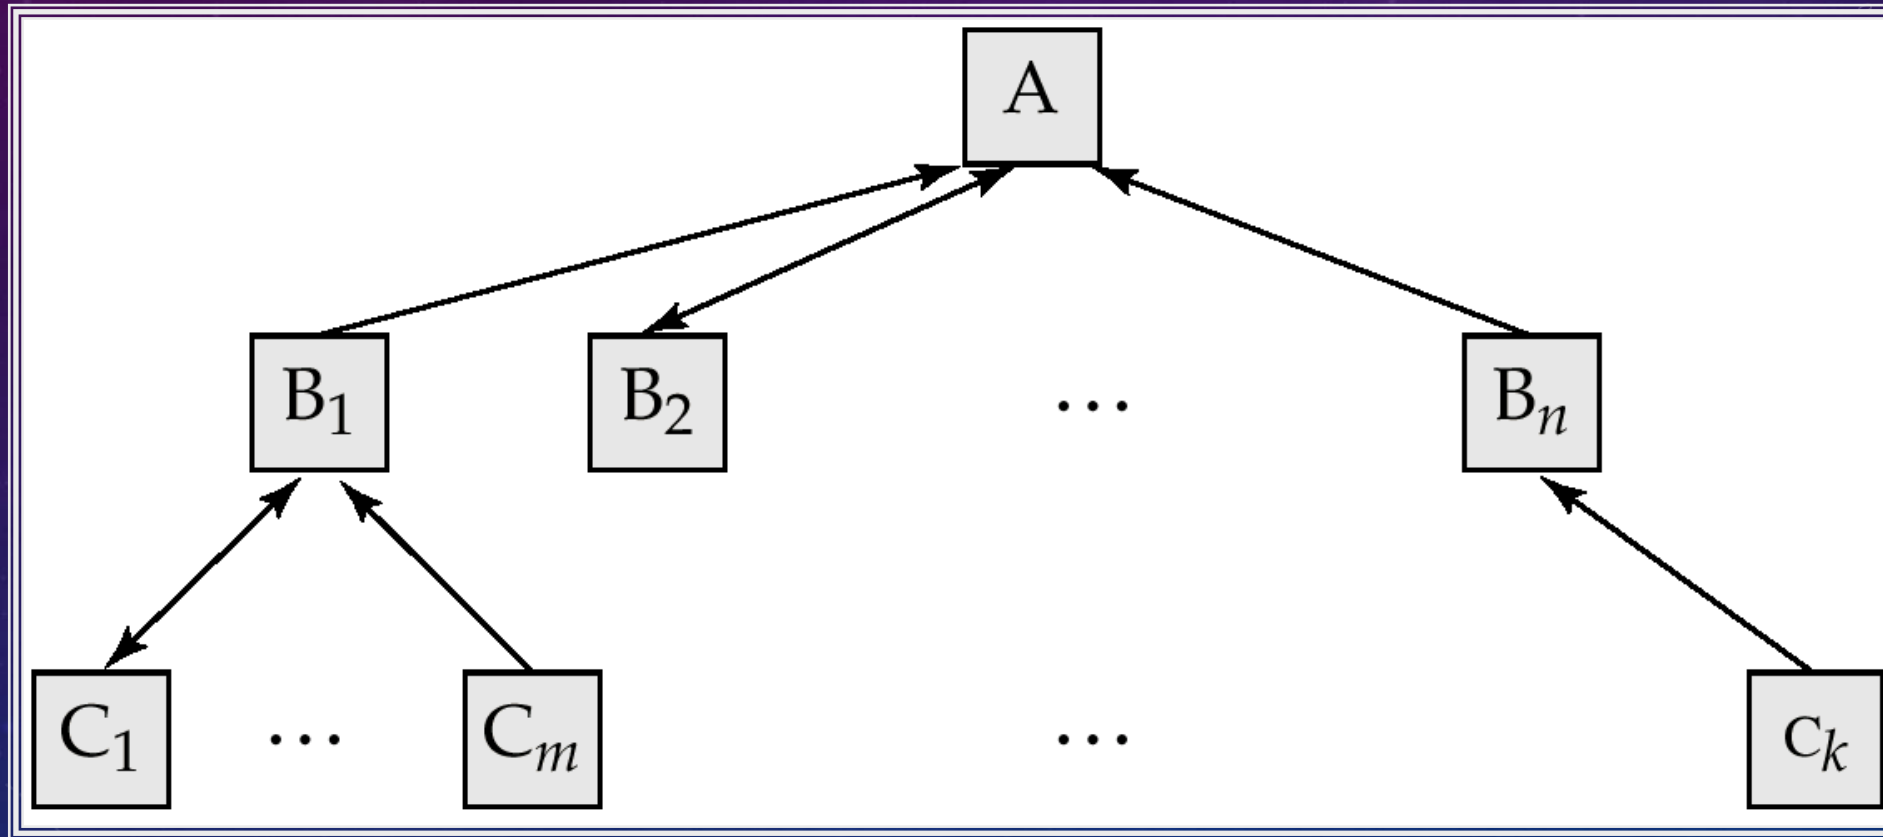
# ENTITY-RELATIONSHIP MODEL

- Entity-relationship model
  - 实体关系模型E-R
  - Entity
  - Relationship
  - Attribute 属性





# HIERARCHICAL DATA MODEL



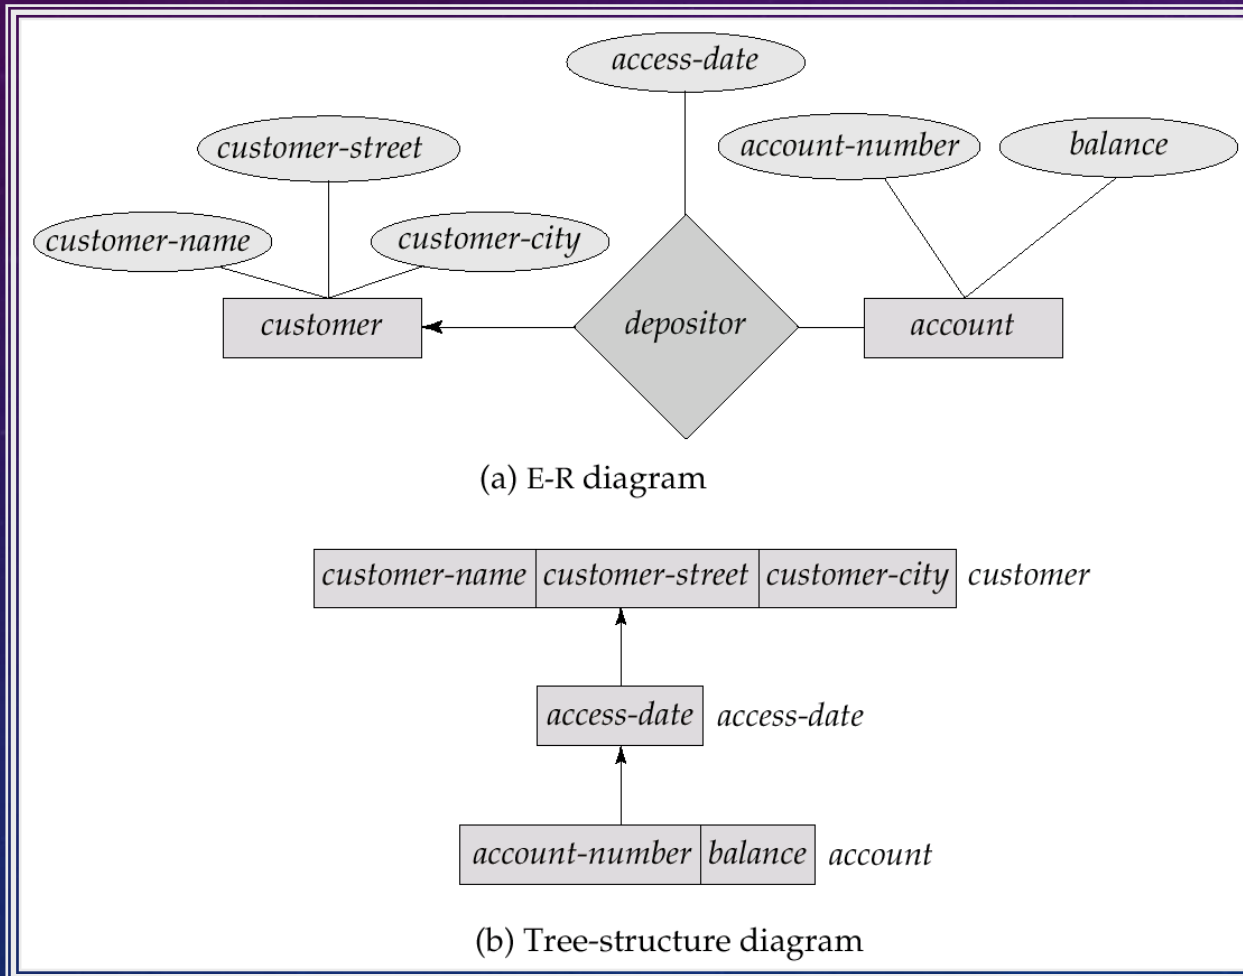
# HIERARCHICAL DATA MODEL

- A hierarchical database consists of a collection of *records* which are connected to one another through *links*.
- a record is a collection of fields, each of which contains only one data value.
- A link is an association between precisely two records.
- The hierarchical model differs from the network model in that the records are organized as collections of trees rather than as arbitrary graphs.

# HIERARCHICAL DATA MODEL

- The schema for a hierarchical database consists of
  - *boxes*, which correspond to record types
  - *lines*, which correspond to links
- Record types are organized in the form of a *rooted tree*.
  - **No cycles** in the underlying graph.
  - Relationships formed in the graph must be such that only one-to-many or one-to-one relationships exist between a parent and a child.

# HIERARCHICAL DATA MODEL: CASE





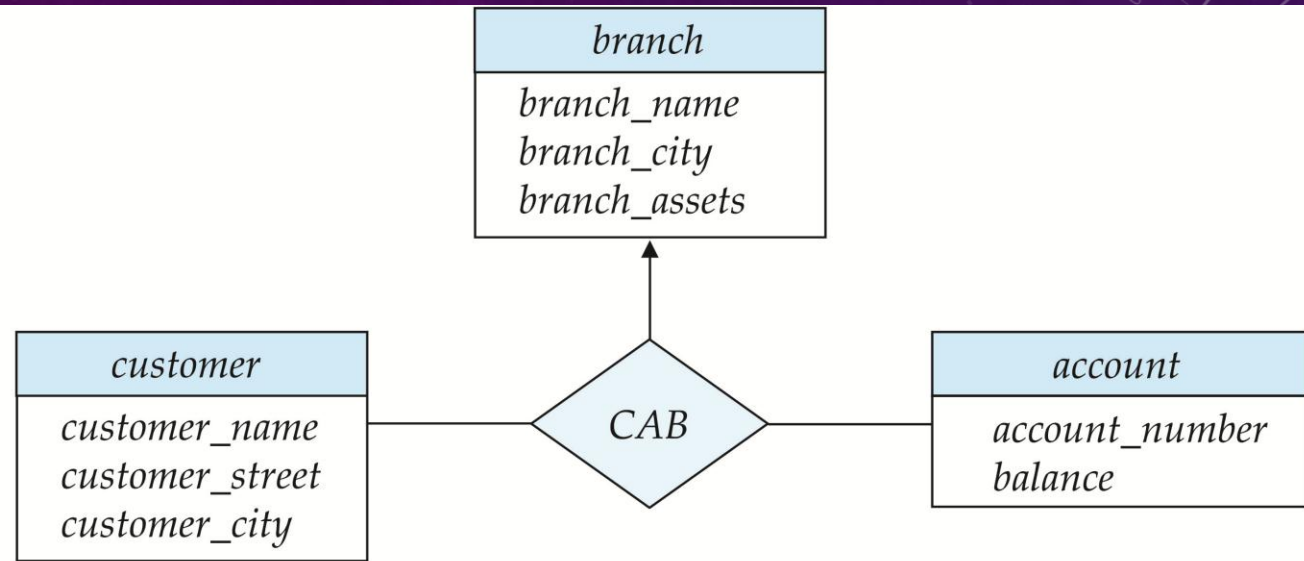
# NETWORK DATA MODEL

- Data are represented by collections of *records*.
  - similar to an entity in the E-R model
  - Records and their fields are represented as *record type*
- Relationships among data are represented by *links*
  - similar to a restricted (binary) form of an E-R relationship
  - restrictions on links depend on whether the relationship is many-many, many-to-one, or one-to-one.

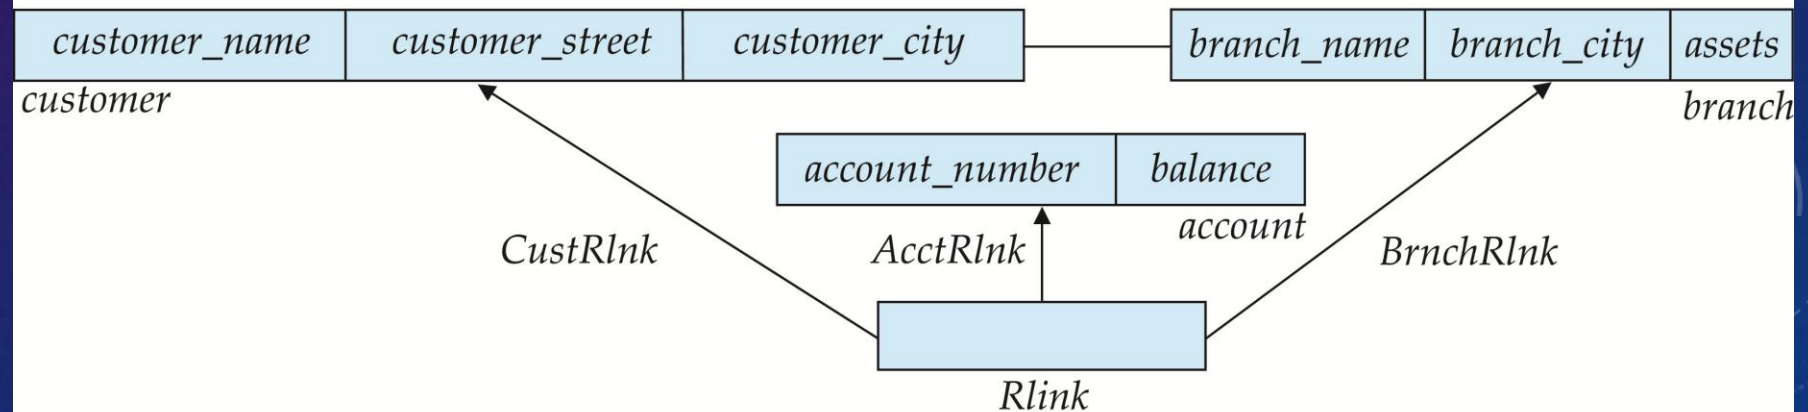
# NETWORK DATA MODEL

- Schema representing the design of a network database.
- A data-structure diagram consists of two basic components:
  - **Boxes**, which correspond to record types.
  - **Lines**, which correspond to links.
- Specifies the overall logical structure of the database.

# NETWORK DATA MODEL : CASE



(a) E-R diagram



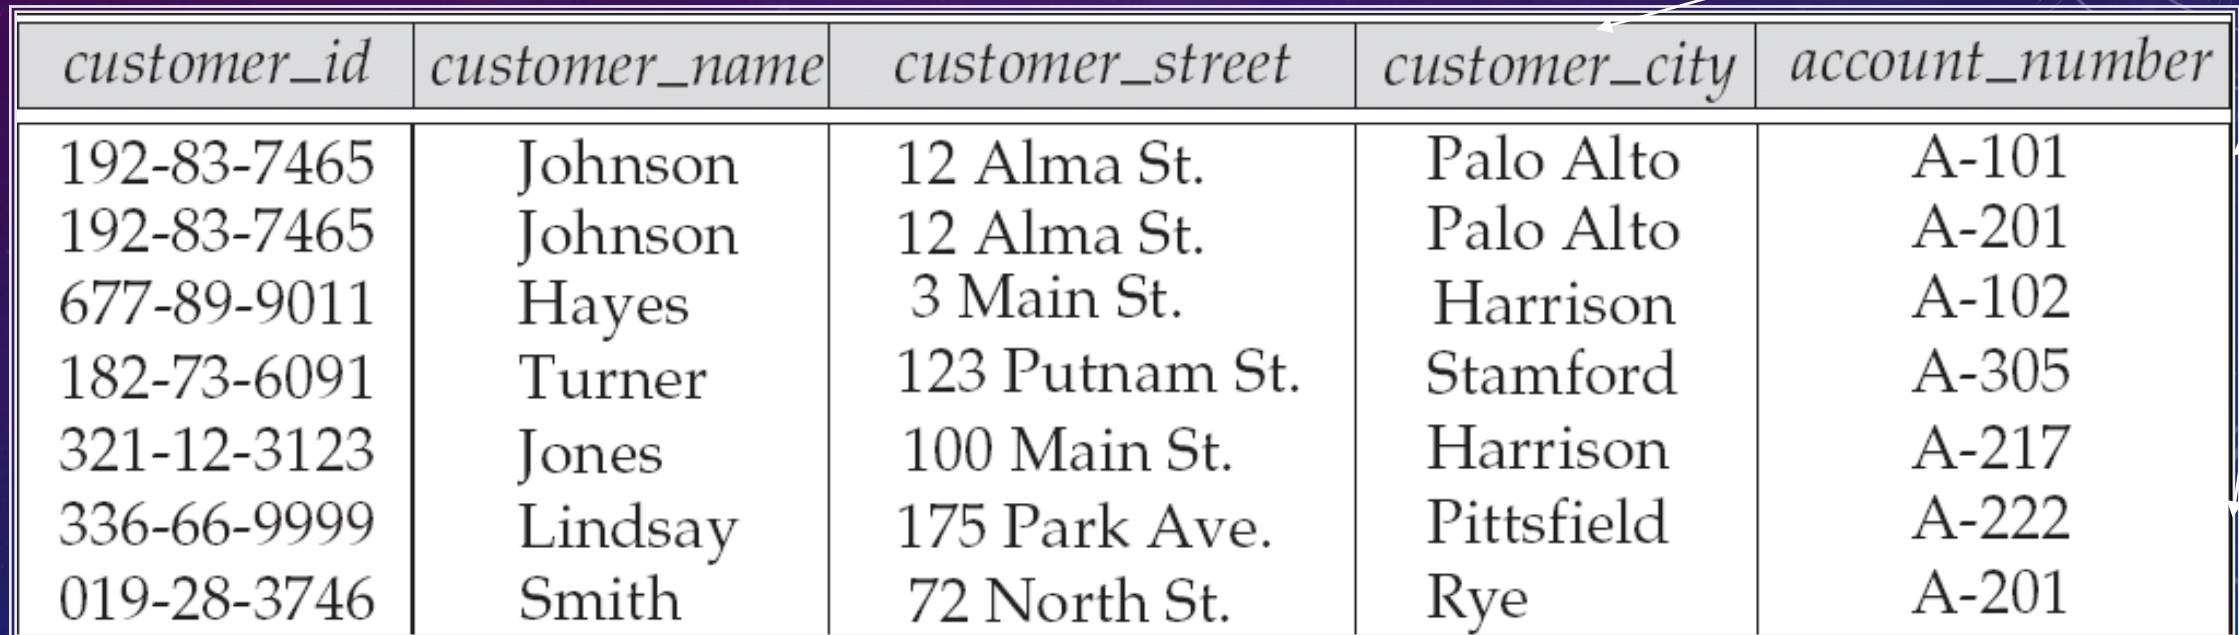
(b) Data structure diagram

# RELATIONAL DATA MODEL

- Relational data model
  - A collection of tables
  - Table
    - A collection of records
    - Multiple columns
  - Relation



# RELATIONAL DATA MODEL: CASE



<i>customer_id</i>	<i>customer_name</i>	<i>customer_street</i>	<i>customer_city</i>	<i>account_number</i>
192-83-7465	Johnson	12 Alma St.	Palo Alto	A-101
192-83-7465	Johnson	12 Alma St.	Palo Alto	A-201
677-89-9011	Hayes	3 Main St.	Harrison	A-102
182-73-6091	Turner	123 Putnam St.	Stamford	A-305
321-12-3123	Jones	100 Main St.	Harrison	A-217
336-66-9999	Lindsay	175 Park Ave.	Pittsfield	A-222
019-28-3746	Smith	72 North St.	Rye	A-201

# RELATIONAL DATA MODEL: CASE

<i>customer-id</i>	<i>customer-name</i>	<i>customer-street</i>	<i>customer-city</i>
192-83-7465	Johnson	12 Alma St.	Palo Alto
019-28-3746	Smith	4 North St.	Rye
677-89-9011	Hayes	3 Main St.	Harrison
182-73-6091	Turner	123 Putnam Ave.	Stamford
321-12-3123	Jones	100 Main St.	Harrison
336-66-9999	Lindsay	175 Park Ave.	Pittsfield
019-28-3746	Smith	72 North St.	Rye

(a) The *customer* table


<i>account-number</i>	<i>balance</i>
A-101	500
A-215	700
A-102	400
A-305	350
A-201	900
A-217	750
A-222	700

(b) The *account* table

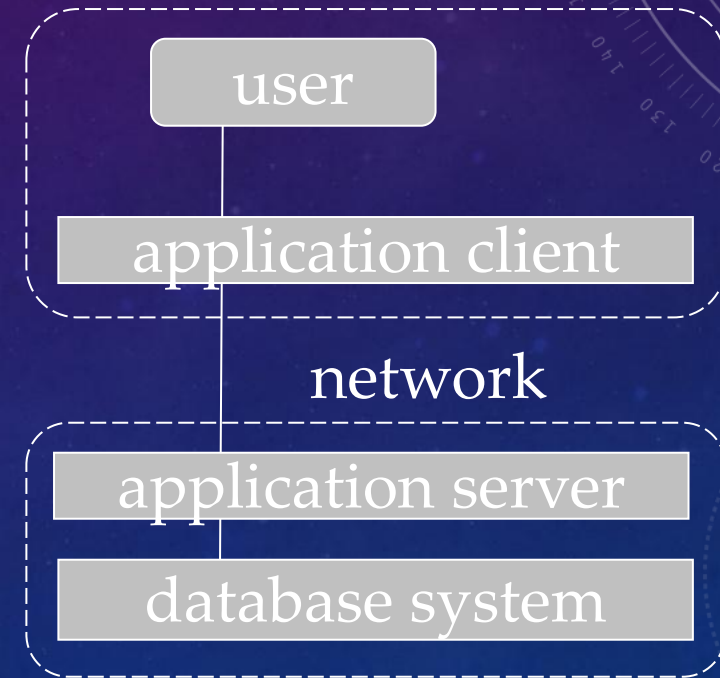
<i>customer-id</i>	<i>account-number</i>
192-83-7465	A-101
192-83-7465	A-201
019-28-3746	A-215
677-89-9011	A-102
182-73-6091	A-305
321-12-3123	A-217
336-66-9999	A-222
019-28-3746	A-201

(c) The *depositor* table

# ARCHITECTURE OF APPLICATION BASED ON CENTRALIZED DATABASE

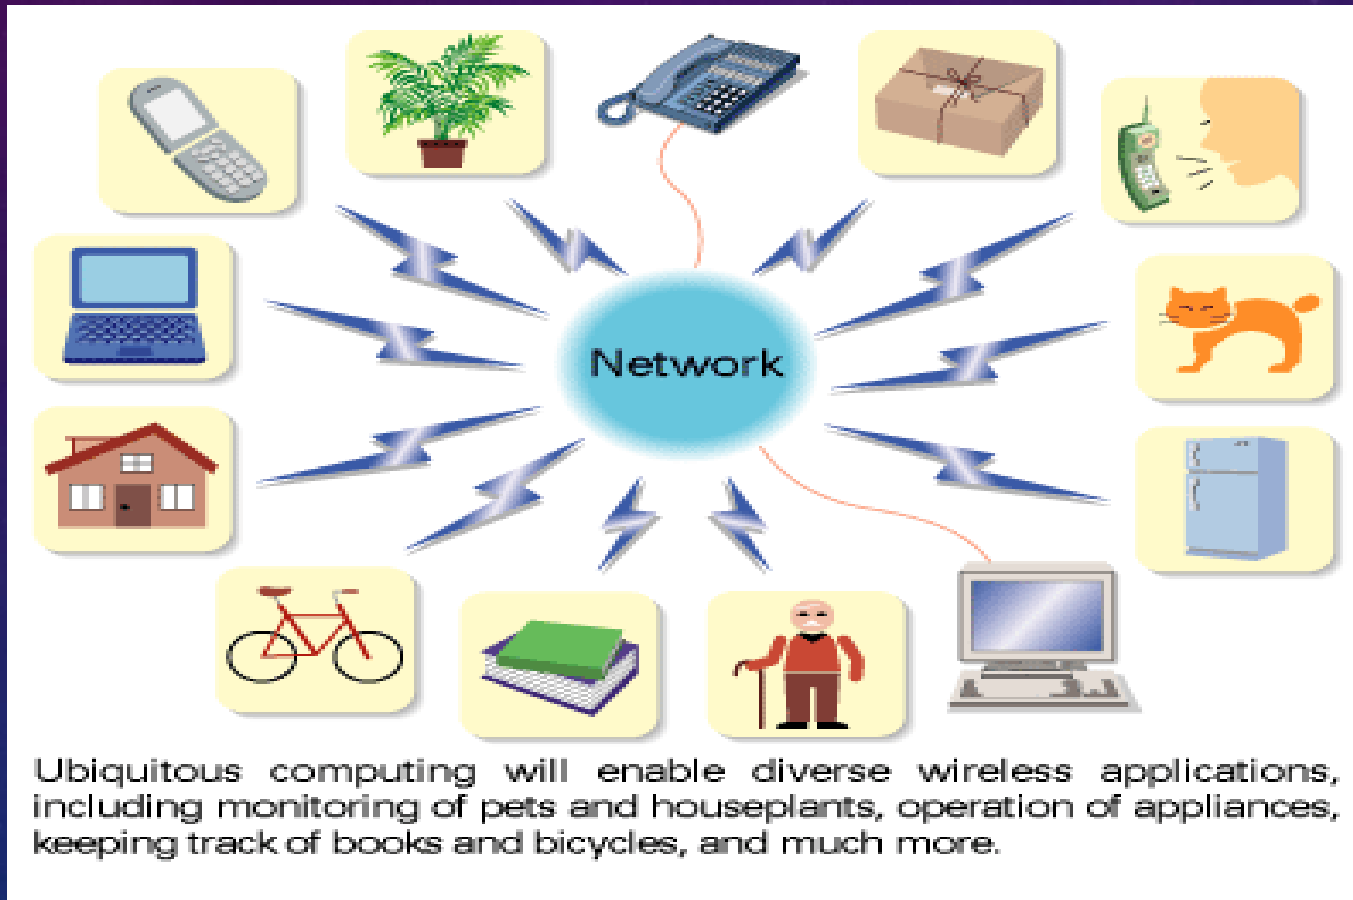


(a) Two-tier  
architecture



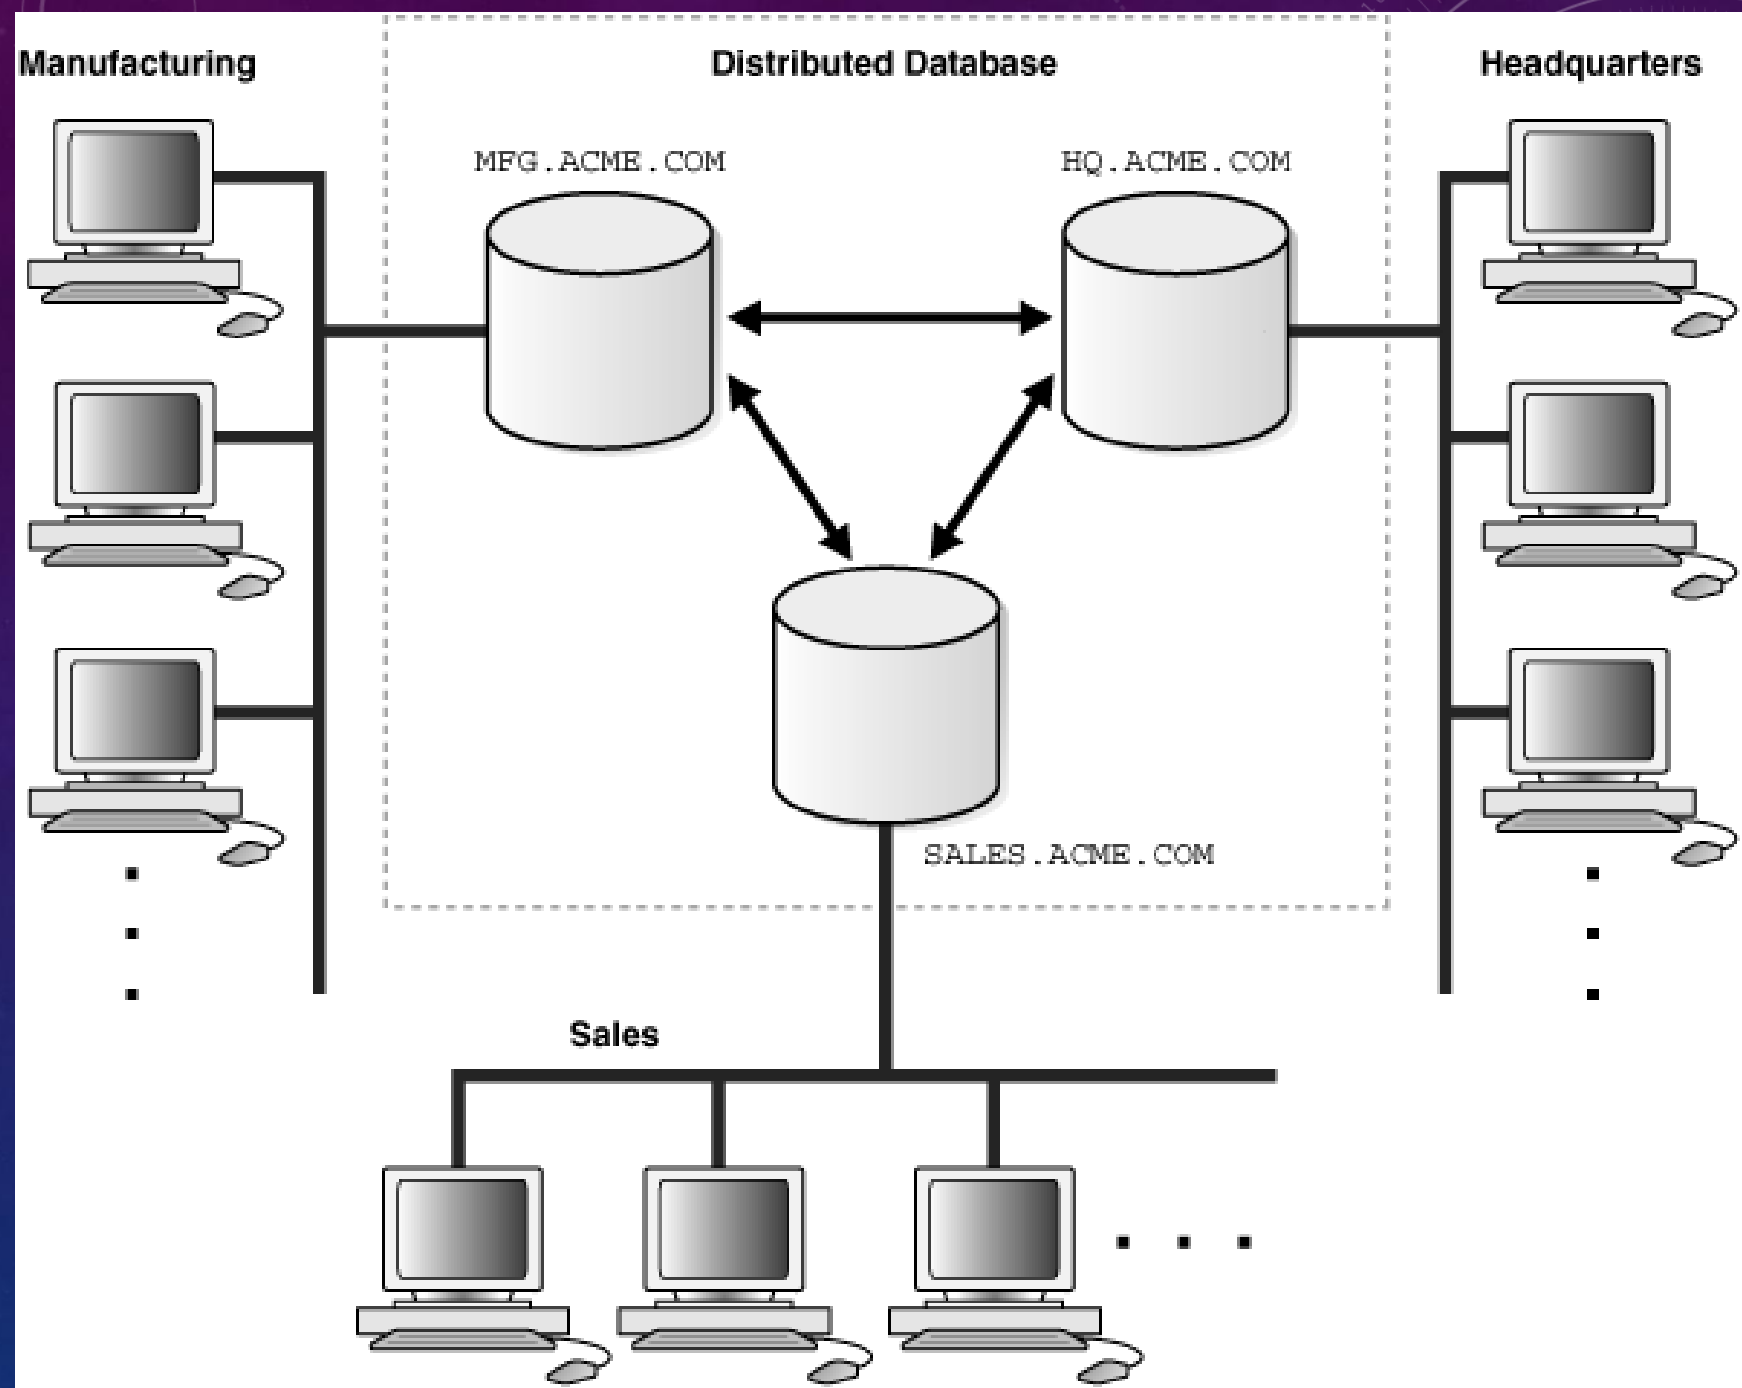
(b) Three-tier  
architecture

# DIGITAL DATA EXPLOSION ...

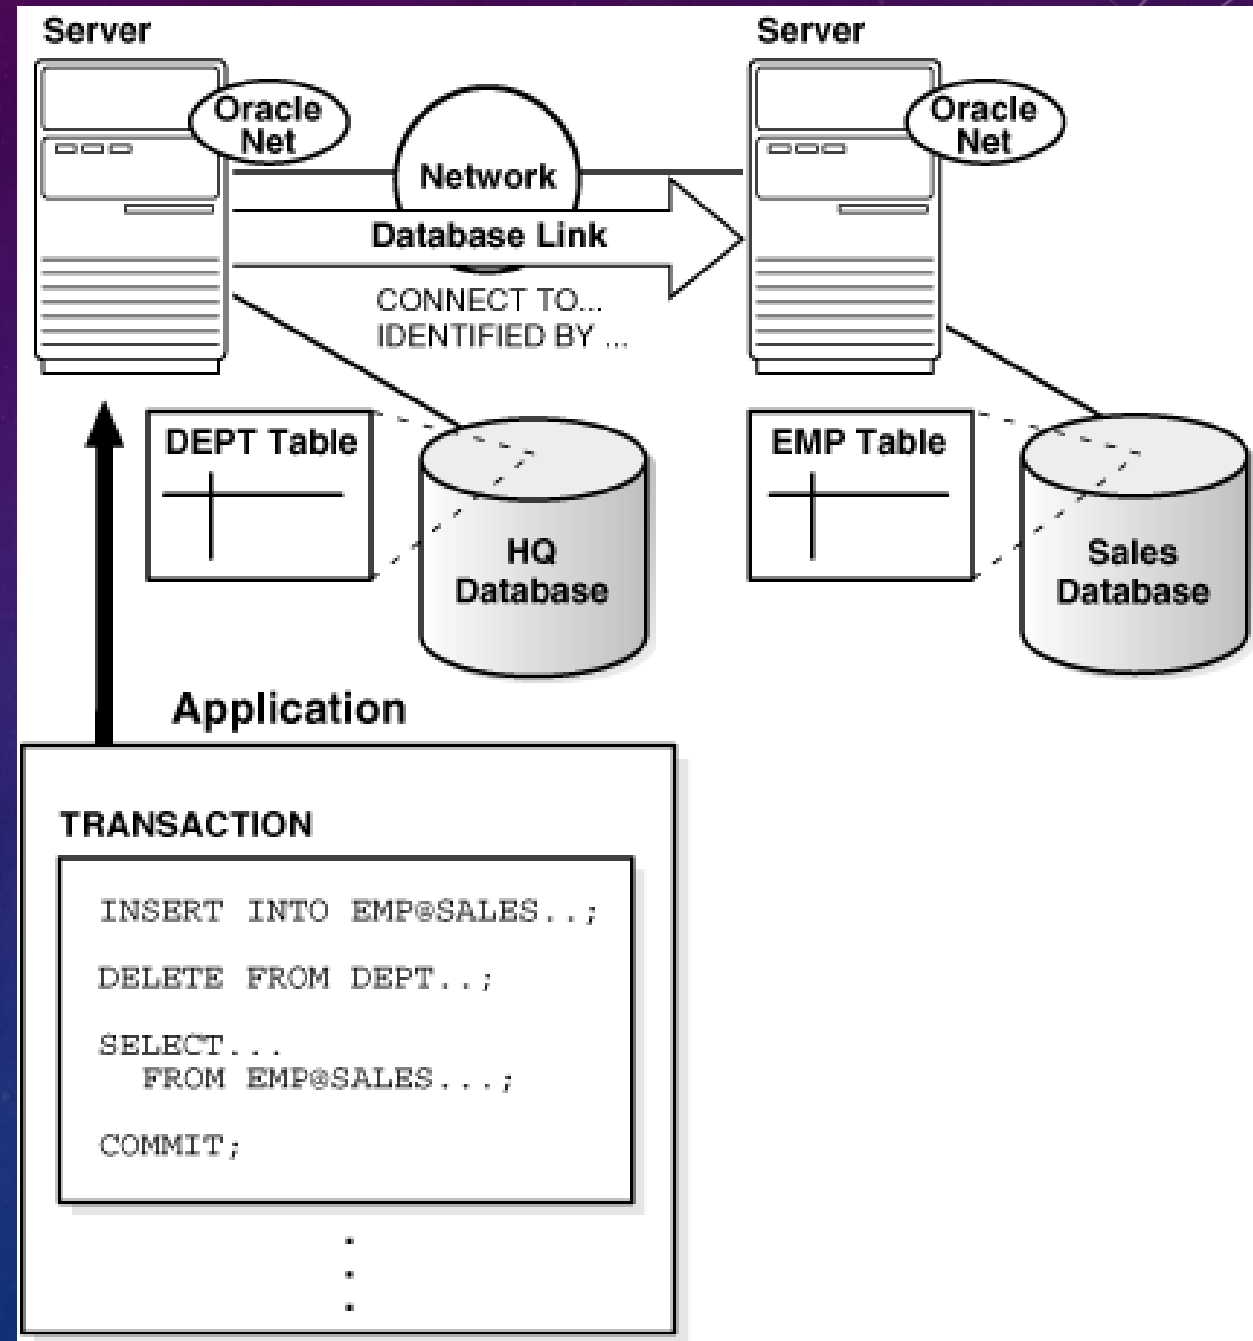




# DISTRIBUTED DATABASE



# DISTRIBUTED DATABASE (ORACLE)



The most important objective  
of the database technology is  
**integration**(集成), not  
**centralization**(集中)

# OBJECTIVES

- The Age of Database
  - Distributed Data Processing
  - What is a Distributed Database System?
  - Data Delivery Alternatives
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD



# DISTRIBUTED DATA PROCESSING

- Distributed processing, Distributed computing
- Distributed computing system
  - It is a number of **autonomous processing elements** (not necessarily homogeneous) that are interconnected by a **computer network** and that **cooperate** in performing their assigned tasks.
  - The “processing element” referred to in this definition is a computing device that can execute a program on its own.

# DISTRIBUTED DATA PROCESSING

## – WHAT IS BEING DISTRIBUTED?

- (1) processing logic处理逻辑 is distributed
  - the definition of a distributed computing system given above implicitly assumes that the **processing logic** or **processing elements** are distributed

# DISTRIBUTED DATA PROCESSING

## – WHAT IS BEING DISTRIBUTED?

- (2) function功能 is distributed
  - Various **functions** of a computer system could be delegated to various pieces of hardware or software



# DISTRIBUTED DATA PROCESSING

## – WHAT IS BEING DISTRIBUTED?

- (3) data数据 is distributed
  - Data used by a number of applications may be distributed to a number of processing sites



# DISTRIBUTED DATA PROCESSING

## – WHAT IS BEING DISTRIBUTED?

- (4) control控制 is distributed
  - The **control** of the execution of various tasks might be distributed instead of being performed by one computer system.

# DISTRIBUTED DATA PROCESSING

## – WHY DO WE DISTRIBUTE AT ALL?

- (1) The classical answers to this question indicate that **distributed processing** **better** corresponds to the organizational structure of today's widely distributed enterprises, and that such a system is more **reliable** and more **responsive**.
- More importantly, many of the current applications of computer technology are inherently distributed.

# DISTRIBUTED DATA PROCESSING

## – WHY DO WE DISTRIBUTE AT ALL?

- (2) it can be stated that the fundamental reason behind distributed processing is to be better able to cope with the large-scale data management problems that we face today, by using a variation of the well-known **divide-and-conquer rule**分而治之.



# OBJECTIVES

- The Age of Database
  - Distributed Data Processing
  - **What is a Distributed Database System?**
  - Data Delivery Alternatives
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD



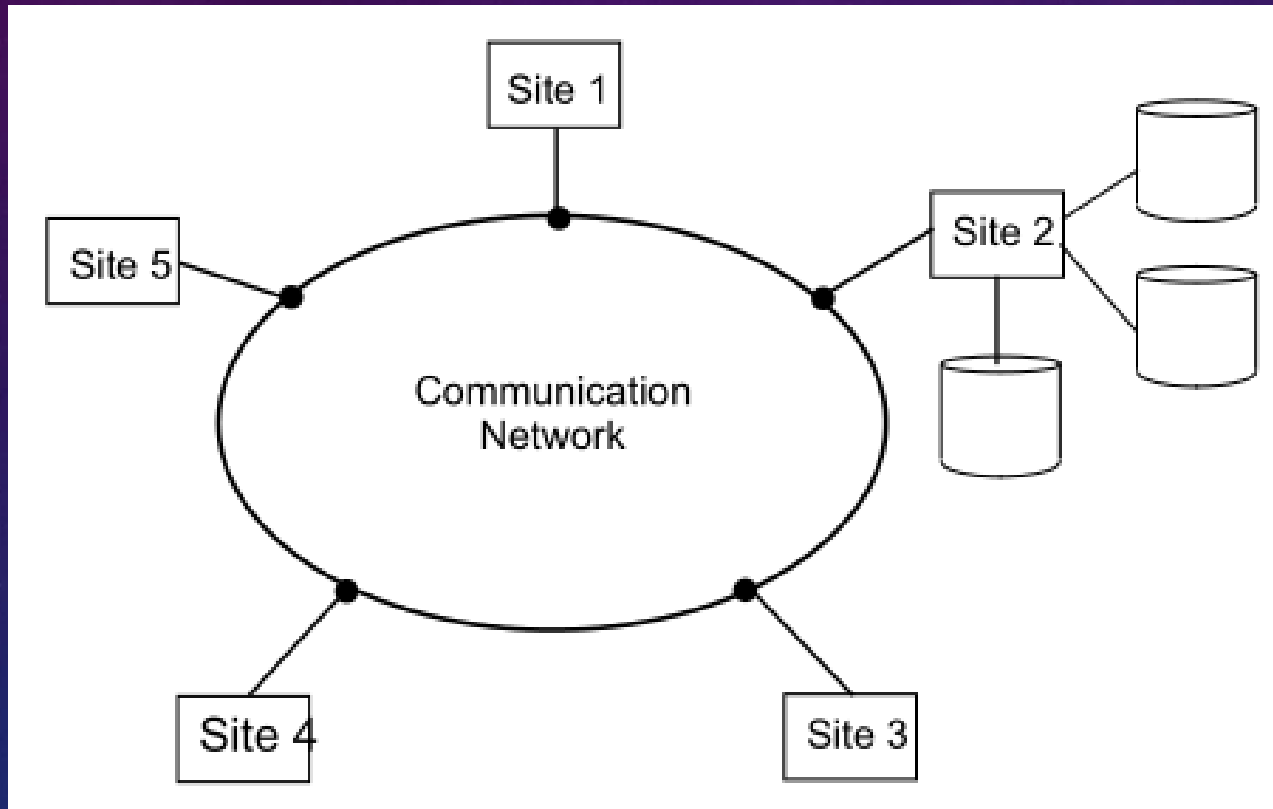
# DISTRIBUTED DATABASE SYSTEM (DDBS)

- Distributed database system (**DDBS**)
  - a collection of multiple, logically interrelated databases distributed over a computer network
    - (1) the distributed database
    - (2) the distributed DBMS

# DISTRIBUTED DATABASE SYSTEM (DDBS)

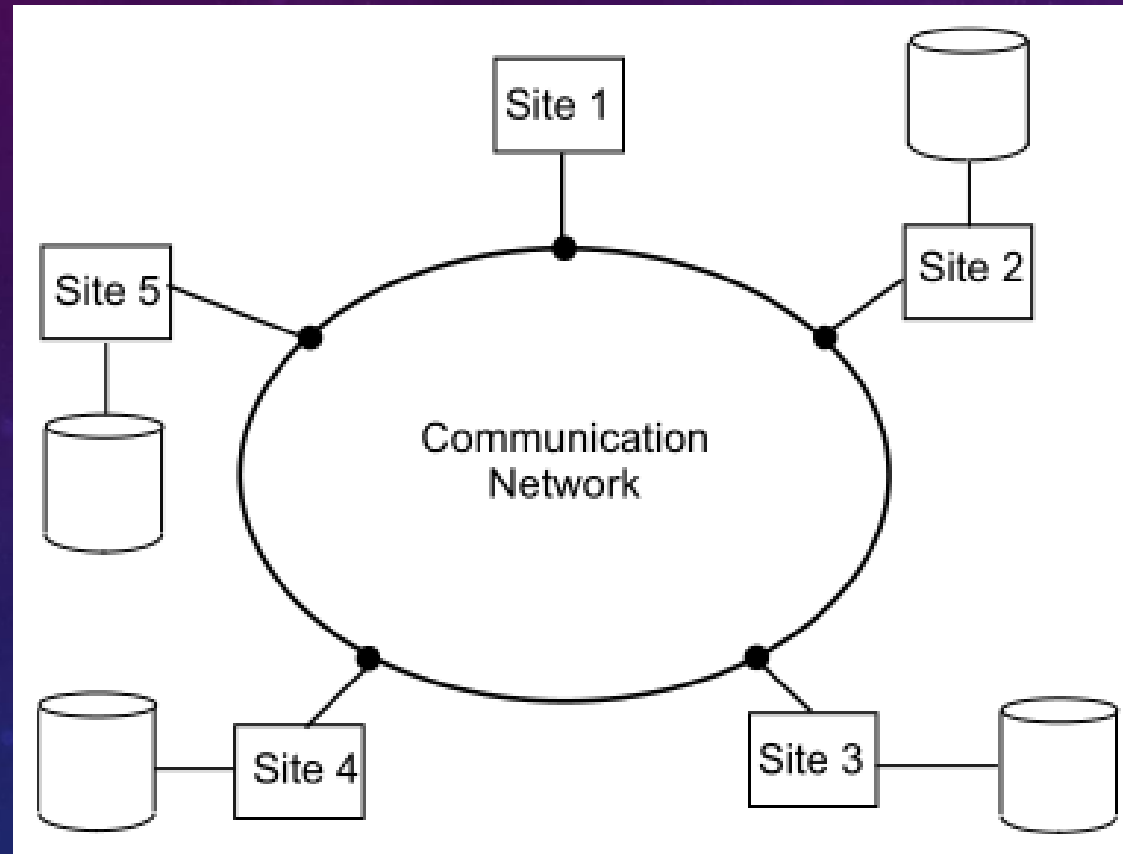
- the **physical distribution of data** is important
  - Note that physical distribution does not necessarily imply that the computer systems be geographically far apart; they could actually be in the same room.
  - It simply implies that the communication between them is done over a network instead of through shared memory or shared disk (as would be the case with multiprocessor systems), with the network as the only shared resource.

# DISTRIBUTED DATABASE SYSTEM (DDBS)



Central Database on a Network

# DISTRIBUTED DATABASE SYSTEM (DDBS)



- logically interrelated
- distributed over a computer network

DDBS Environment



# OBJECTIVES

- The Age of Database
  - Distributed Data Processing
  - What is a Distributed Database System?
  - **Data Delivery Alternatives**
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD

# DATA DELIVERY ALTERNATIVES

- three orthogonal dimensions:
  - delivery modes
    - Pull-only, push-only, hybrid
  - Frequency
    - Periodic, conditional, ad-hoc凭经验的, irregular
  - communication methods
    - Unicast单播, one-to-many

# OBJECTIVES

- The Age of Database
  - Distributed Data Processing
  - What is a Distributed Database System?
  - Data Delivery Alternatives
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD



# PROMISES承诺 OF DDBSS

- Transparent Management of Distributed and Replicated Data
- Reliability Through Distributed Transactions
- Improved Performance
- Easier System Expansion



# PROMISES承诺 OF DDBSS

- **Transparent Management** of Distributed and Replicated Data
- Reliability Through Distributed Transactions
- Improved Performance
- Easier System Expansion

## PROMISES承诺 OF DDBSS

### – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Transparency refers to separation of the higher-level semantics of a system from lower-level implementation issues. In other words, a transparent system “hides” the implementation details from users.
- The advantage of a fully transparent DBMS is the high level of support that it provides for the development of complex applications.
- It is obvious that we would like to make all DBMSs (centralized or distributed) fully transparent.

# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Example:

- EMP( ENO , ENAME, TITLE)
- PROJ( PNO , PNAME, BUDGET)
- SAL( TITLE, AMT )
- ASG( ENO, PNO , RESP, DUR)

```
SELECT ENAME, AMT  
FROM EMP, ASG, SAL  
WHERE ASG.DUR > 12  
AND EMP.ENO = ASG.ENO  
AND SAL.TITLE = EMP.TITLE
```



# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Example (cont.):

business, performance and reliability

- **Fragmentation**分片

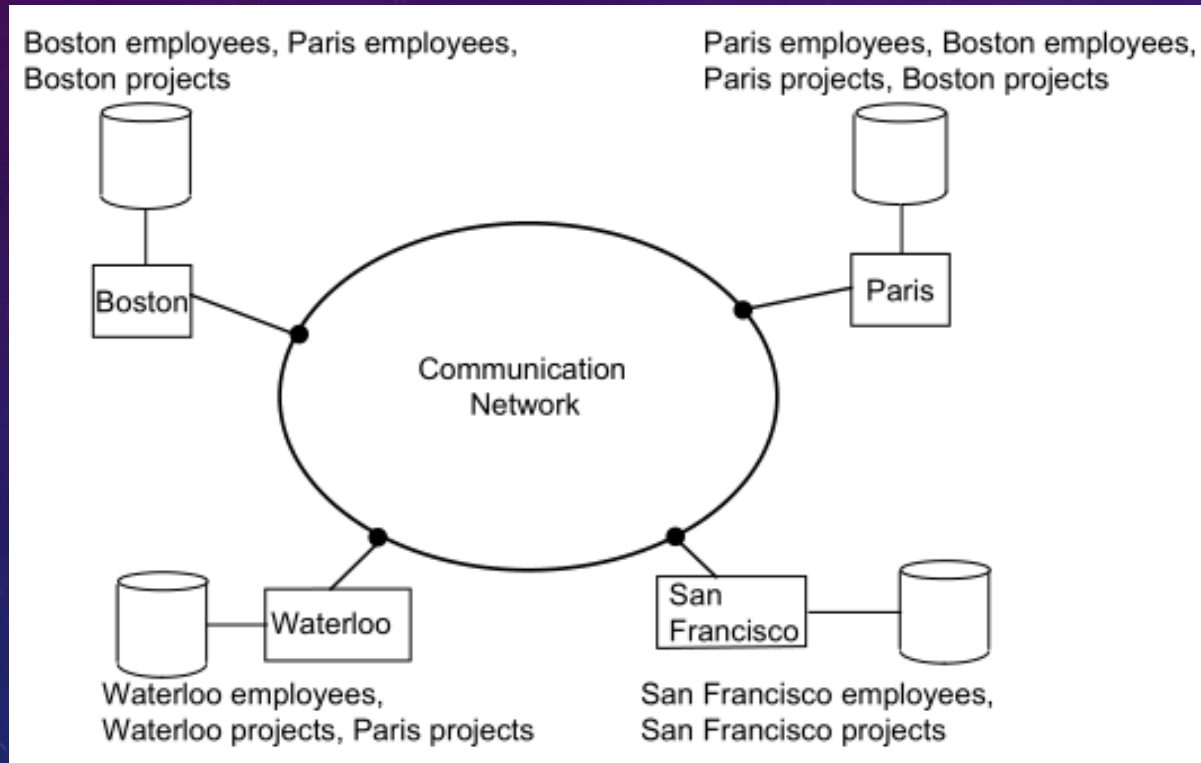
- the employees in Waterloo office are stored in **Waterloo**, those in the Boston office are stored in **Boston**, and so forth.

```
SELECT ENAME, AMT  
FROM EMP, ASG, SAL  
WHERE ASG.DUR > 12  
AND EMP.ENO = ASG.ENO  
AND SAL.TITLE = EMP.TITLE
```



# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA



performance and reliability

**Transparent** operations

Under  
fragmentation and replication data

**A Distributed Application**

# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- types of transparencies:\*
  - Data Independence
  - Network Transparency
  - Replication Transparency
  - Fragmentation Transparency

# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Data Independence
  - Schema definition模式定义
    - Logical data independence逻辑数据独立性
  - Physical data description物理数据描述
    - Physical data independence物理数据独立性



# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Network Transparency
  - the user should be protected from the operational details of the network; possibly even hiding the existence of the network.
  - (1) network transparency or distribution transparency
  - (2) location transparency and naming transparency



# PROMISES 承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Replication Transparency

- Multi-copies: let us just mention that for performance, reliability, and availability reasons, it is usually desirable to be able to distribute data in a replicated fashion across the machines on a network.
- From a user's perspective, it is preferable not to be involved with handling copies and having to specify the fact that a certain action can and/or should be taken on multiple copies.

# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Fragmentation Transparency
  - it is commonly desirable to divide each database relation into smaller fragments and treat each fragment as a separate database object (i.e., another relation).
  - (1)Horizontal fragmentation水平分片
  - (2)Vertical fragmentation垂直分片
  - a translation from what is called a **global query** to several **fragment queries**

# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Fragmentation Transparency
  - (1) Horizontal fragmentation水平分片
    - a relation is partitioned into a set of sub-relations each of which have a subset of the tuples (rows) of the original relation



# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Fragmentation Transparency
  - (2) Vertical fragmentation垂直分片
    - each sub-relation is defined on a subset of the attributes (columns) of the original relation



# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Who Should Provide Transparency?
  - A full transparency makes the management of distributed data very difficult
  - the level of transparency is inevitably a compromise折衷 between ease of use and the difficulty and overhead cost of providing high levels of transparency


# PROMISES承诺 OF DDBSS

## – TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA

- Who Should Provide Transparency: three layers
  - (1) We could leave the responsibility of providing transparent access to data resources to the access layer
  - (2) The second layer at which transparency can be provided is the operating system level
  - (3) The third layer at which transparency can be supported is within the DBMS

# PROMISES 承诺 OF DDBSS

– TRANSPARENT MANAGEMENT OF DISTRIBUTED AND REPLICATED DATA



# PROMISES承诺 OF DDBSS

- Transparent Management of Distributed and Replicated Data
- **Reliability** Through Distributed Transactions
- Improved Performance
- Easier System Expansion



# PROMISES承诺 OF DDBSS

## – RELIABILITY THROUGH DISTRIBUTED TRANSACTIONS

- single points of failure单点故障
- Transactions and transaction processing
  - A transaction is a basic unit of **consistent** and **reliable** computing, consisting of a sequence of database operations executed as an atomic action

# PROMISES承诺 OF DDBSS

## – RELIABILITY THROUGH DISTRIBUTED TRANSACTIONS

- Concurrency transparency并发透明
  - It transforms a consistent database state to another consistent database state even when a number of such transactions are executed concurrently, and even when failures occur (also called failure atomicity)
- Failure atomicity故障原子性

# PROMISES承诺 OF DDBSS

- Transparent Management of Distributed and Replicated Data
- Reliability Through Distributed Transactions
- Improved Performance
- Easier System Expansion



# PROMISES 承诺 OF DDBSS

## – IMPROVED PERFORMANCE



- (1) data localization 数据本地化
  - a distributed DBMS fragments the conceptual database, enabling data to be stored in close proximity to its points of use
- advantages:
  - Since each site handles only a portion of the database, contention for CPU and I/O services is not as severe as for centralized databases.
  - Localization reduces remote access delays that are usually involved in wide area networks



# PROMISES承诺 OF DDBSS

## – IMPROVED PERFORMANCE

- (2) the inherent parallelism
  - the inherent parallelism of distributed systems may be exploited for inter-query and intra-query parallelism
  - it results from the ability to execute multiple queries at the same time while intra-query parallelism is achieved by breaking up a single query into a number of sub queries each of which is executed at a different site, accessing a different part of the distributed database.

# PROMISES承诺 OF DDBSS

- Transparent Management of Distributed and Replicated Data
- Reliability Through Distributed Transactions
- Improved Performance
- **Easier System Expansion**

# PROMISES承诺 OF DDBSS

## – EASIER SYSTEM EXPANSION

- In a distributed environment, it is much easier to accommodate increasing database sizes
- Grosh's law
  - it was commonly believed that it would be possible to purchase a fourfold powerful computer if one spent twice as much



# OBJECTIVES

- The Age of Database
  - Distributed Data Processing
  - What is a Distributed Database System?
  - Data Delivery Alternatives
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD



# COMPLICATIONS INTRODUCED BY DISTRIBUTION

- (1) data may be replicated in a distributed environment.
  - A distributed database can be designed so that the entire database, or portions of it, reside at different sites of a computer network.

# COMPLICATIONS INTRODUCED BY DISTRIBUTION

- (2)effect and durability
  - if some sites fail (e.g., by either hardware or software malfunction),
  - or if some communication links fail (making some of the sites unreachable) while an update is being executed,
  - the system must make sure that the effects will be reflected on the data residing at the failing or unreachable sites as soon as the system can recover from the failure.

# COMPLICATIONS INTRODUCED BY DISTRIBUTION

- (3) distributed transactions and synchronization
  - since each site cannot have instantaneous information on the actions currently being carried out at the other sites,
  - the synchronization of transactions on multiple sites is considerably harder than for a centralized system.

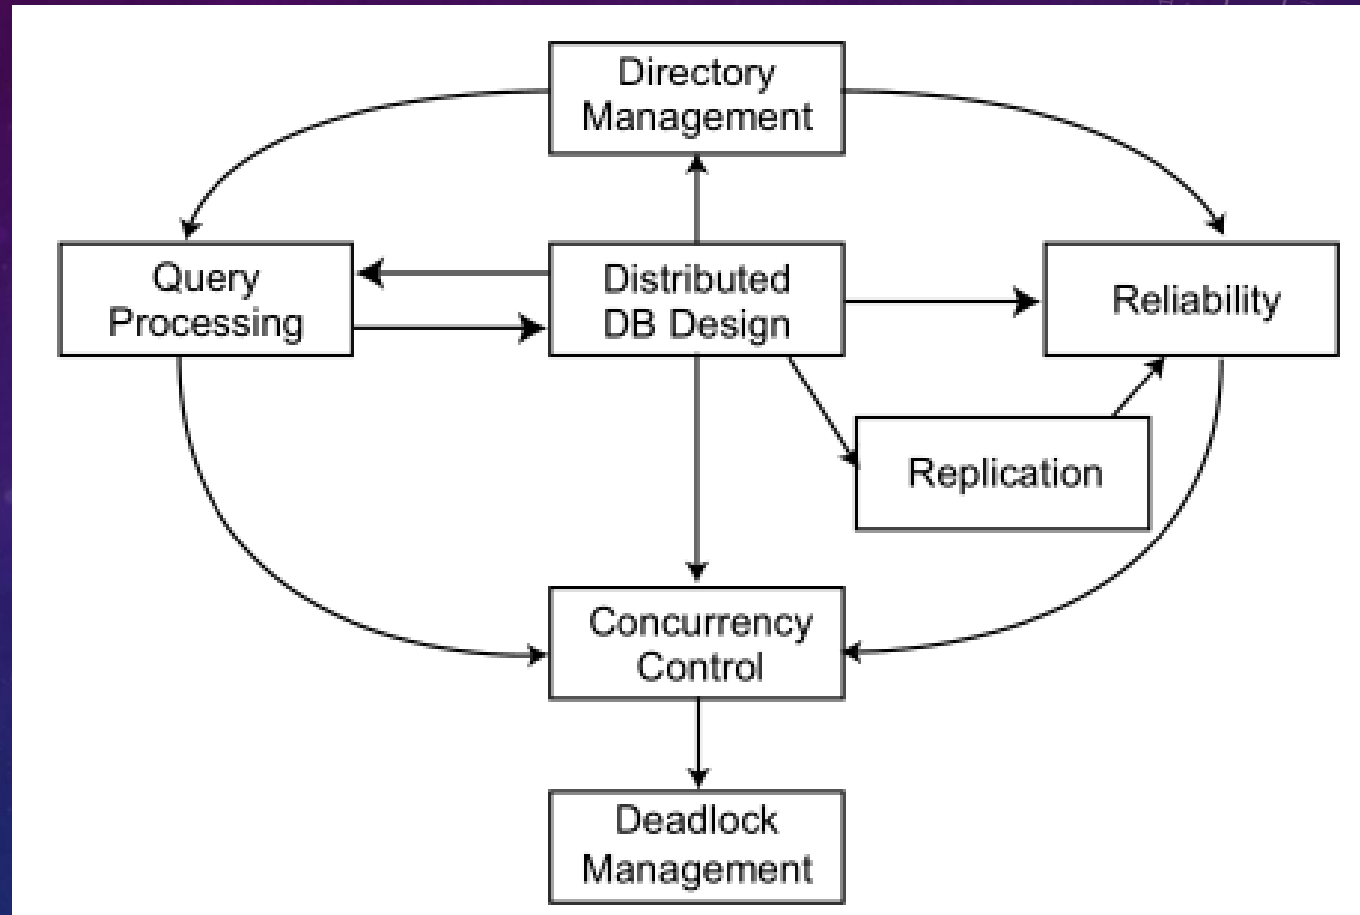


# OBJECTIVES

- The Age of Database
  - Distributed Data Processing
  - What is a Distributed Database System?
  - Data Delivery Alternatives
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD



# DESIGN ISSUES



Relationship Among Research Issues

# DESIGN ISSUES

- (1) Distributed Database Design
- two basic alternatives to placing data
  - (a) partitioned (or non-replicated)划分
  - (b) replicated复制
    - Fully replicated, partially replicated

# DESIGN ISSUES

- (1) Distributed Database Design
- two fundamental design issues:
  - Fragmentation(分片,划分): the separation of the database into partitions called fragments 片段
  - Distribution(分配,分布): the optimum distribution of fragments
- The research in this area mostly involves mathematical programming in order to minimize the combined cost of storing the database, processing transactions against it, and message communication among sites

# DESIGN ISSUES

- (2) Distributed Directory Management
  - A directory contains information (such as descriptions and locations) about data items in the database.
  - A directory may be **global** to the entire DDBS or **local** to each site;
  - it can be centralized at one site or distributed over several sites; there can be a single copy or multiple copies.



# DESIGN ISSUES

- (3) Distributed Query Processing
  - Query processing deals with designing **algorithms** that analyze queries and convert them into a series of data manipulation operations.
  - The problem is how to decide on a strategy for executing each query over the network in the most cost-effective way, however cost is defined.
  - The factors to be considered are the distribution of data, communication costs, and lack of sufficient locally-available information.
  - The objective is to optimize where the inherent parallelism is used to improve the performance of executing the transaction, subject to the above-mentioned constraints.
  - NP-hard

# DESIGN ISSUES

- (4) Distributed Concurrency Control 并发控制
  - **Concurrency control** involves the synchronization of accesses to the distributed database, such that the integrity 完整性 of the database is maintained.
  - It is one of the most extensively studied problems in the DDBS field
  - This problem in a distributed context is somewhat different than in a centralized framework. One not only has to worry about the integrity of a single database, but also about the consistency of multiple copies of the database
  - mutual consistency 相互一致性: The condition that requires all the values of multiple copies of every data item to converge to the same value

# DESIGN ISSUES

- (4) Distributed Concurrency Control (cont.,)
- mutual consistency相互一致性:
  - The condition that requires all the values of multiple copies of every data item to converge to the same value
  - (a) pessimistic悲观方法:synchronizing the execution of user requests before the execution starts (**locking** method)
  - (b) optimistic乐观方法:executing the requests and then checking if the execution has compromised the consistency of the database (**timestamping** method)



# DESIGN ISSUES

- (5) Distributed Deadlock Management
  - Deadlock: The competition among users for access to a set of resources can result in a deadlock if the synchronization mechanism is based on locking
  - prevention, avoidance, and detection/recovery
  - Similar to OS deadlock



# DESIGN ISSUES

- (6) Reliability of Distributed DBMS
  - one of the potential advantages of distributed systems is improved **reliability** and **availability**.
  - however, is not a feature that comes automatically.
  - It is important that mechanisms be provided to ensure the consistency of the database as well as to detect failures and recover from them

# DESIGN ISSUES

- (7) Replication

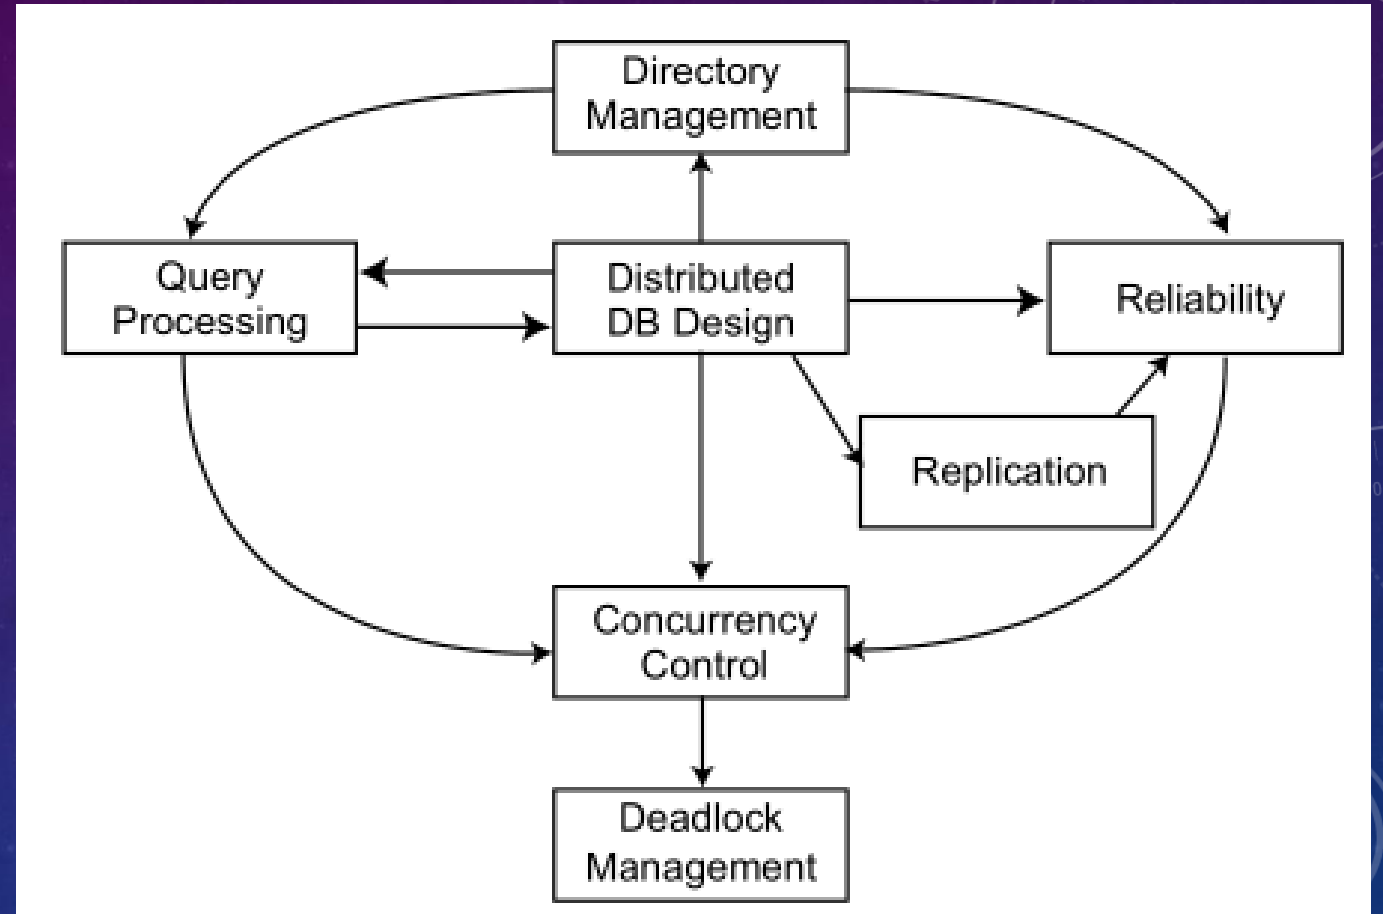
- If the distributed database is (**partially** or **fully**) replicated, it is necessary to implement **protocols** that ensure the **consistency** of the replicas, i.e., copies of the same data item have the same value.
- (a) These protocols can be **eager** 及时 in that they force the updates to be applied to all the replicas before the transaction completes
- (b) they may be lazy 惰性 so that the transaction updates one copy (called the master 主拷贝) from which updates are **propagated** 传播 to the others after the transaction completes.

# DESIGN ISSUES

## – ADDITIONAL ISSUES

- Multidatabase systems
  - federated databases 联邦数据库
  - data integration systems 数据集成系统
- P2P DB
- Web DB

# DESIGN ISSUES – RELATIONSHIP AMONG PROBLEMS



Relationship Among Research Issues



# OBJECTIVES

- The Age of Database
  - Distributed Data Processing
  - What is a Distributed Database System?
  - Data Delivery Alternatives
  - Promises of DDBSs
  - Complications Introduced by Distribution
  - Design Issues
  - Distributed DBMS Architecture
  - History
- ©LXD

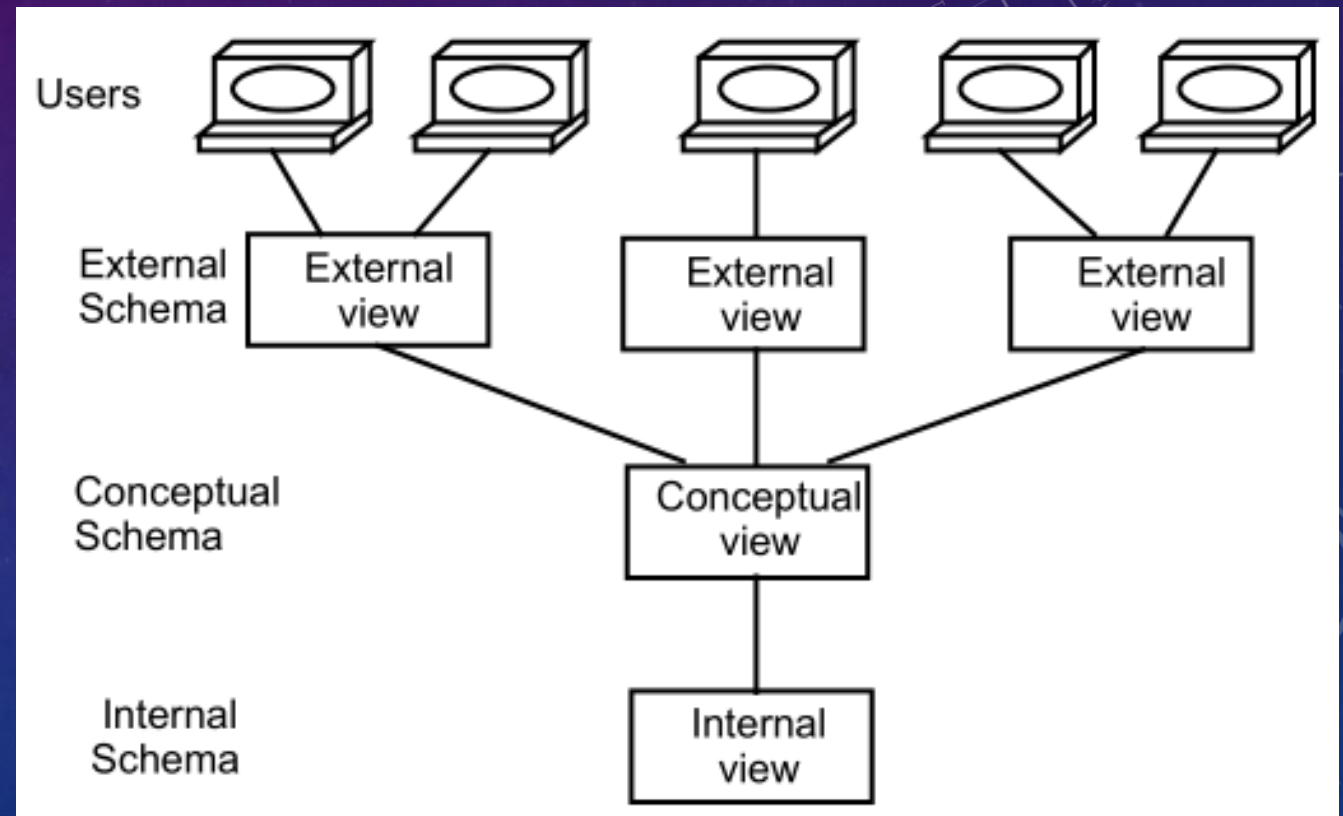
# DISTRIBUTED DBMS ARCHITECTURE

- The architecture of a system defines its structure.
  - This means that the components of the system are identified, the function of each component is specified, and the interrelationships and interactions among these components are defined.
- three “reference” architectures:
  - client/server systems
  - peer-to-peer distributed DBMS
  - multidatabase systems

# DISTRIBUTED DBMS ARCHITECTURE

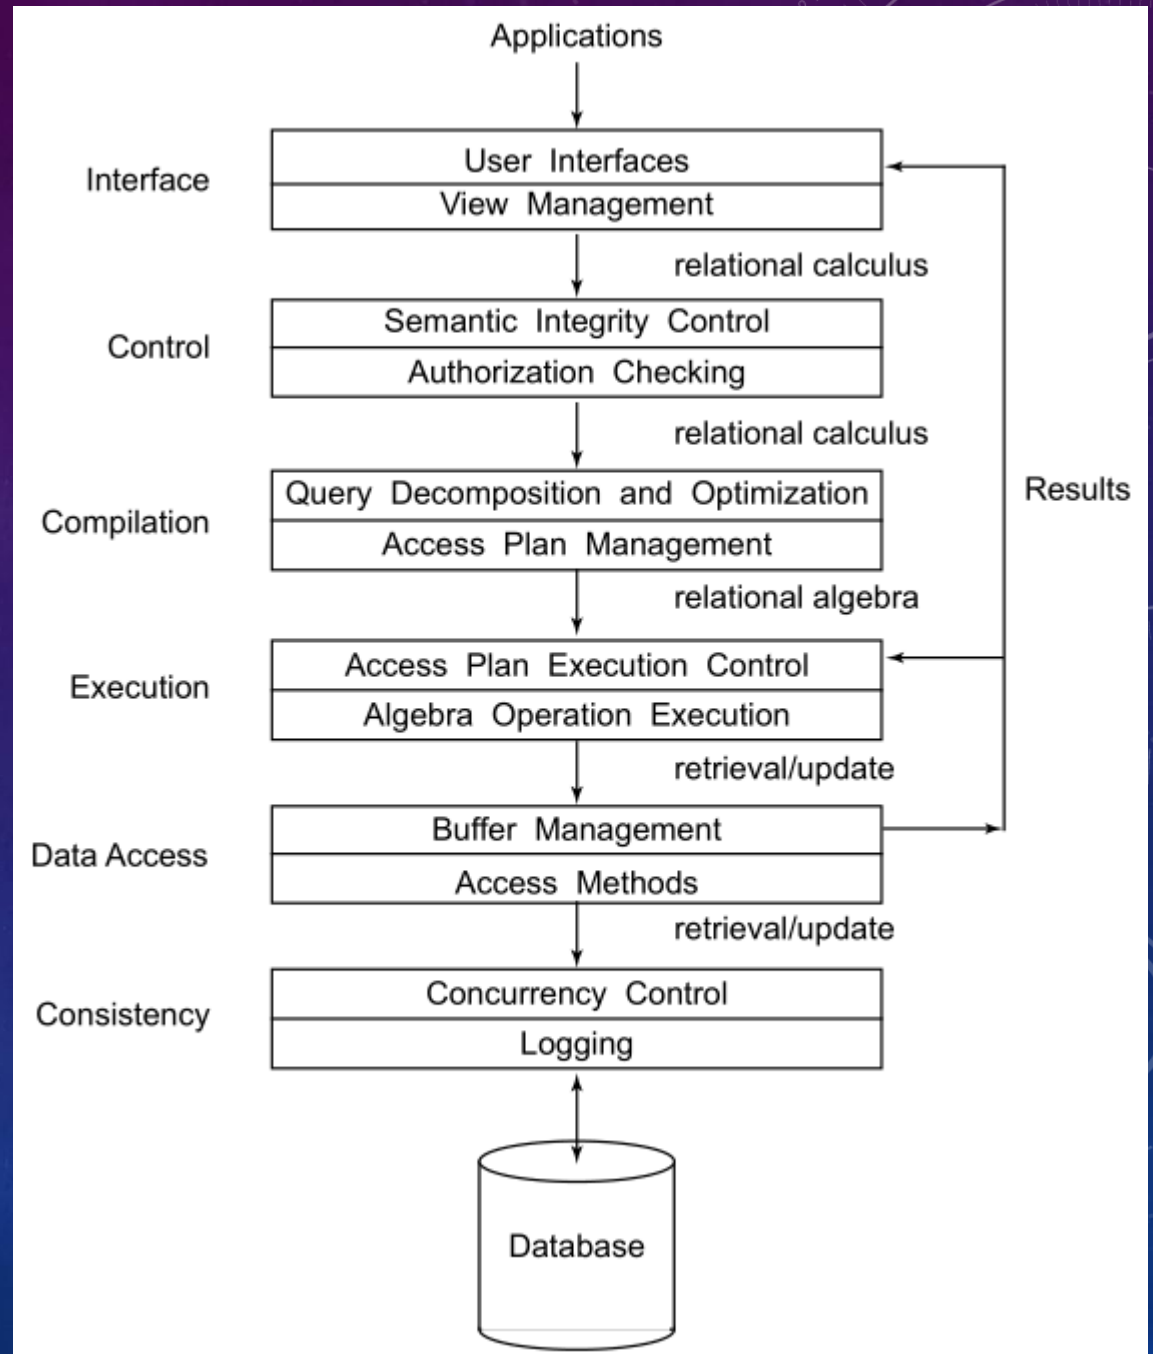
## – ANSI/SPARC ARCHITECTURE

- The ANSI/SPARC Architecture, 1975
  - Internal Schema
  - Conceptual Schema
  - External Schema
- **Data independence**
  - Logical data independence
  - Physical data independence





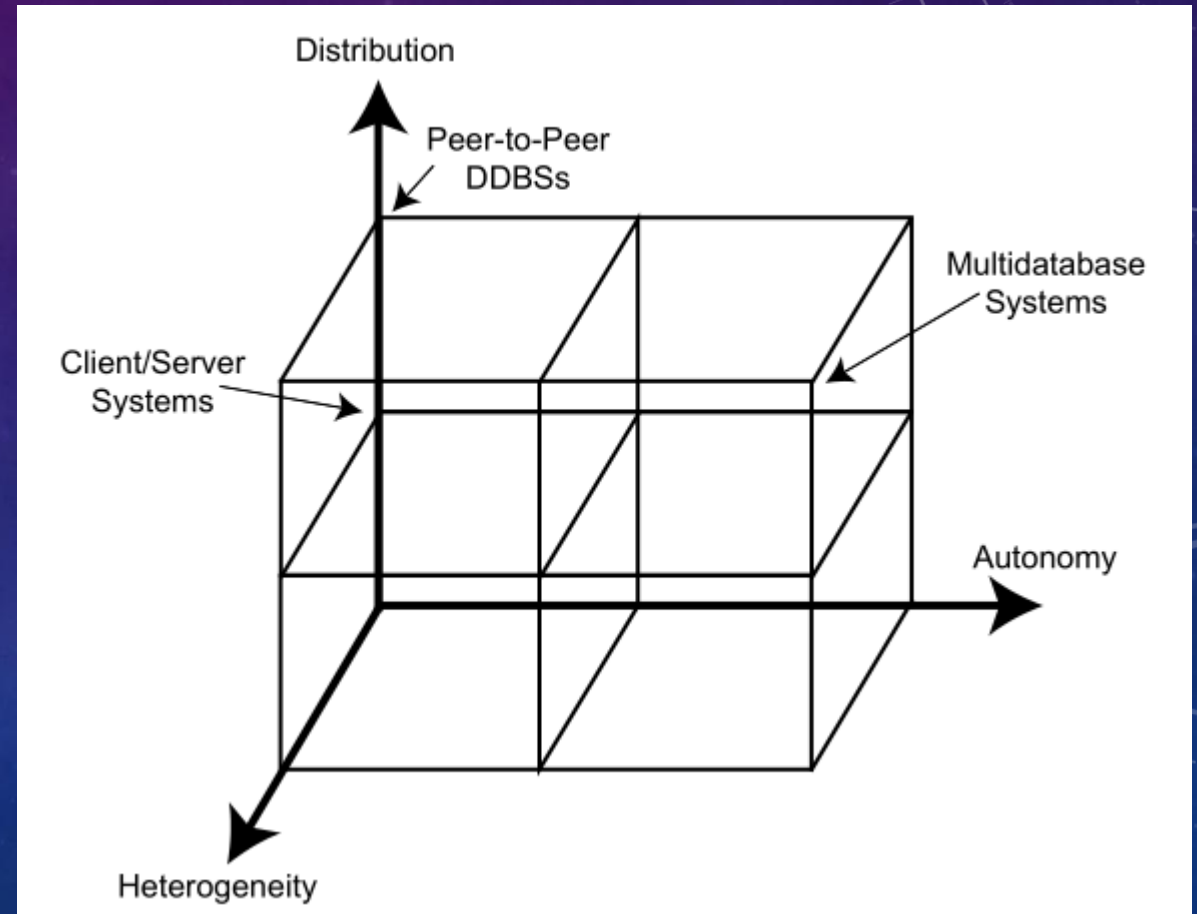
# DISTRIBUTED DBMS ARCHITECTURE – A GENERIC CENTRALIZED DBMS ARCHITECTURE



# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED DBMSS


- classification:
  - (1) the autonomy of local systems
  - (2) their distribution
  - (3) their heterogeneity



# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED DBMSS

- classification:
  - (1) the autonomy of local systems
  - (2) their distribution
  - (3) their heterogeneity



# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED DBMSS

- **Autonomy** 自治性
  - Autonomy, in this context, refers to the distribution of control, not of data.
  - It indicates the degree to which individual DBMSs can operate independently



# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED

- DBMSS  
Autonomy自治性(cont.,)
- Requirements of an autonomous system[Gligor and PopescuZeletin,1986]:
  - (1)The local operations of the individual DBMSs are not affected by their participation in the distributed system.
  - (2)The manner in which the individual DBMSs process queries and optimize them should not be affected by the execution of global queries that access multiple databases.
  - (3)System consistency or operation should not be **compromised**妥协 when individual DBMSs join or leave the distributed system.

# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED

- DBMSS 自治性(cont.,)
- the dimensions of autonomy [Du and Elmagarmid, 1989]:
  - 1. **Design autonomy**: Individual DBMSs are **free** to use the **data models** and **transaction management** techniques that they prefer.
  - 2. **Communication autonomy**: Each of the individual DBMSs is free to **make its own decision** as to what type of information it wants to provide to the other DBMSs or to the software that controls their global execution.
  - 3. **Execution autonomy**: Each DBMS can **execute the transactions** that are submitted to it in any way that it wants to.

# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED

### DBMSS 自治性(cont.,)


- the types of DDBS based on different autonomy
  - tight integration systems 紧密集成的系统
    - where a single-image of the entire database is available to any user who wants to share the information, which may reside in multiple databases.
  - semiautonomous systems 半自治系统
    - it consist of DBMSs that can (and usually do) operate independently, but have decided to participate in a federation to make their local data sharable.
  - total isolation systems 全孤立系统
    - the individual systems are stand-alone DBMSs that know neither of the existence of other DBMSs nor how to communicate with them.



# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED DBMSS

- classification:
  - (1) the autonomy of local systems
  - (2) **their distribution**
  - (3) their heterogeneity



# DISTRIBUTED DBMS ARCHITECTURE

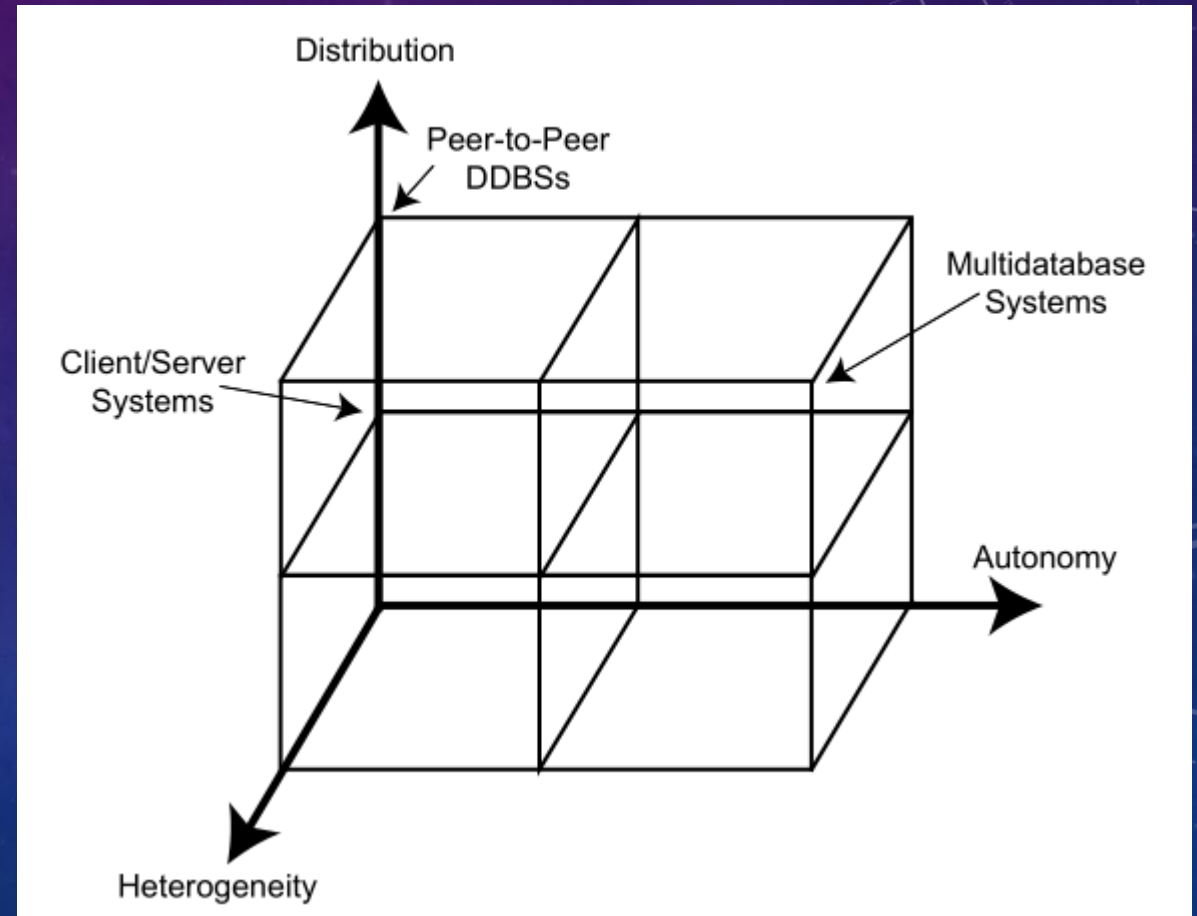
## – ARCHITECTURAL MODELS FOR DISTRIBUTED DBMSS

- **Distribution**分布
  - the physical distribution of data over multiple sites
- **taxonomy**
  - (1) client/server distribution
  - (2) peer-to-peer distribution (or full distribution)
  - (3) non-distributed

# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED DBMSS

- classification:
  - (1) the autonomy of local systems
  - (2) their distribution
  - (3) their heterogeneity





# DISTRIBUTED DBMS ARCHITECTURE

## – ARCHITECTURAL MODELS FOR DISTRIBUTED DBMSS

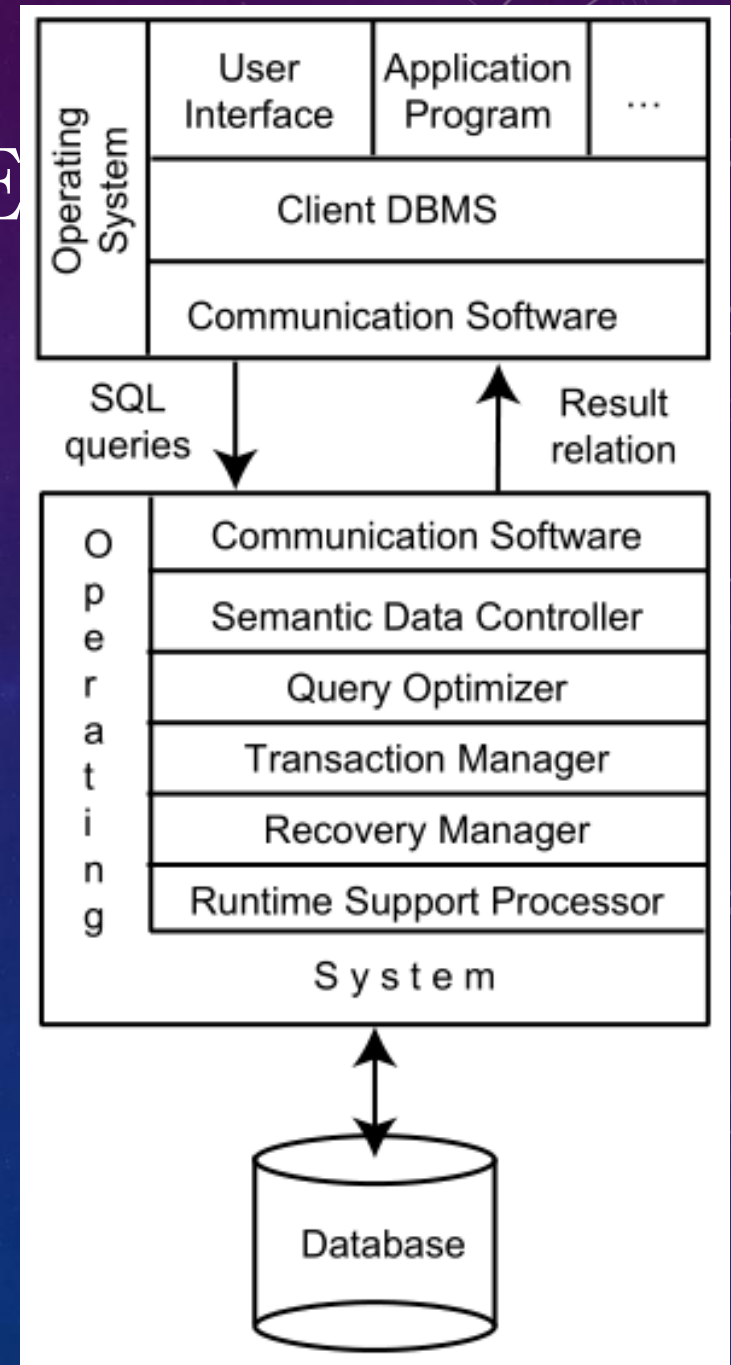
- **Heterogeneity**异构性
  - Heterogeneity may occur in various forms in distributed systems, ranging from hardware heterogeneity and differences in networking protocols to variations in data managers
  - data models, query languages, and transaction management protocols

# DISTRIBUTED DBMS ARCHITECTURAL ALTERNATIVES

# DISTRIBUTED DBMS ARCHITECTURE

## – CLIENT/SERVER SYSTEMS

- client/server computing
- client/server DBMS
  - Multiple client/Single server
  - Multiple client/Multiple server

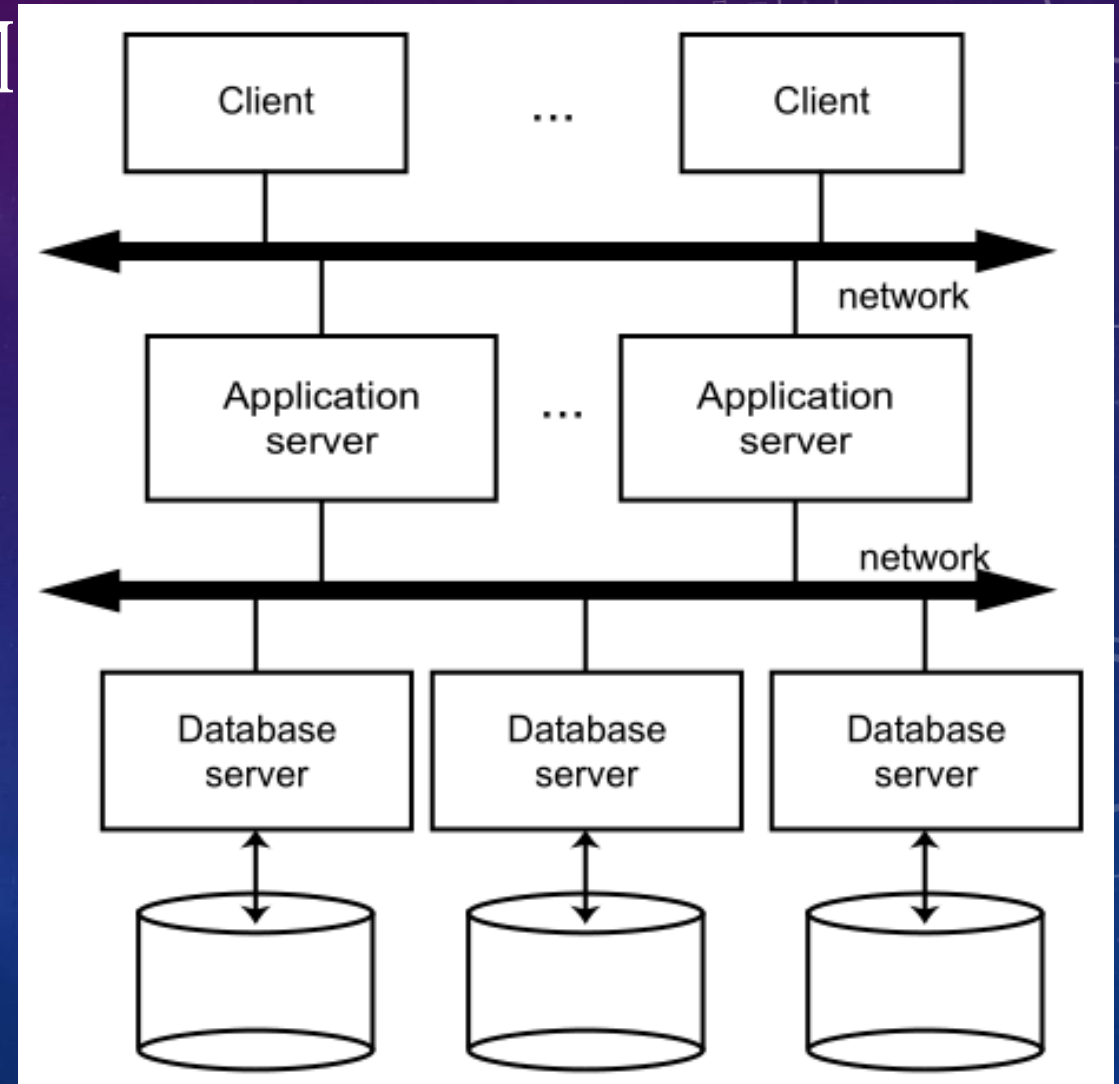




# DISTRIBUTED DBMS ARCHITECTURE

## – CLIENT/SERVER SYSTEM


- C/S/S
  - Client server (web server)
  - Application server
  - Database server



# DISTRIBUTED DBMS ARCHITECTURE

## – PEER-TO-PEER SYSTEMS

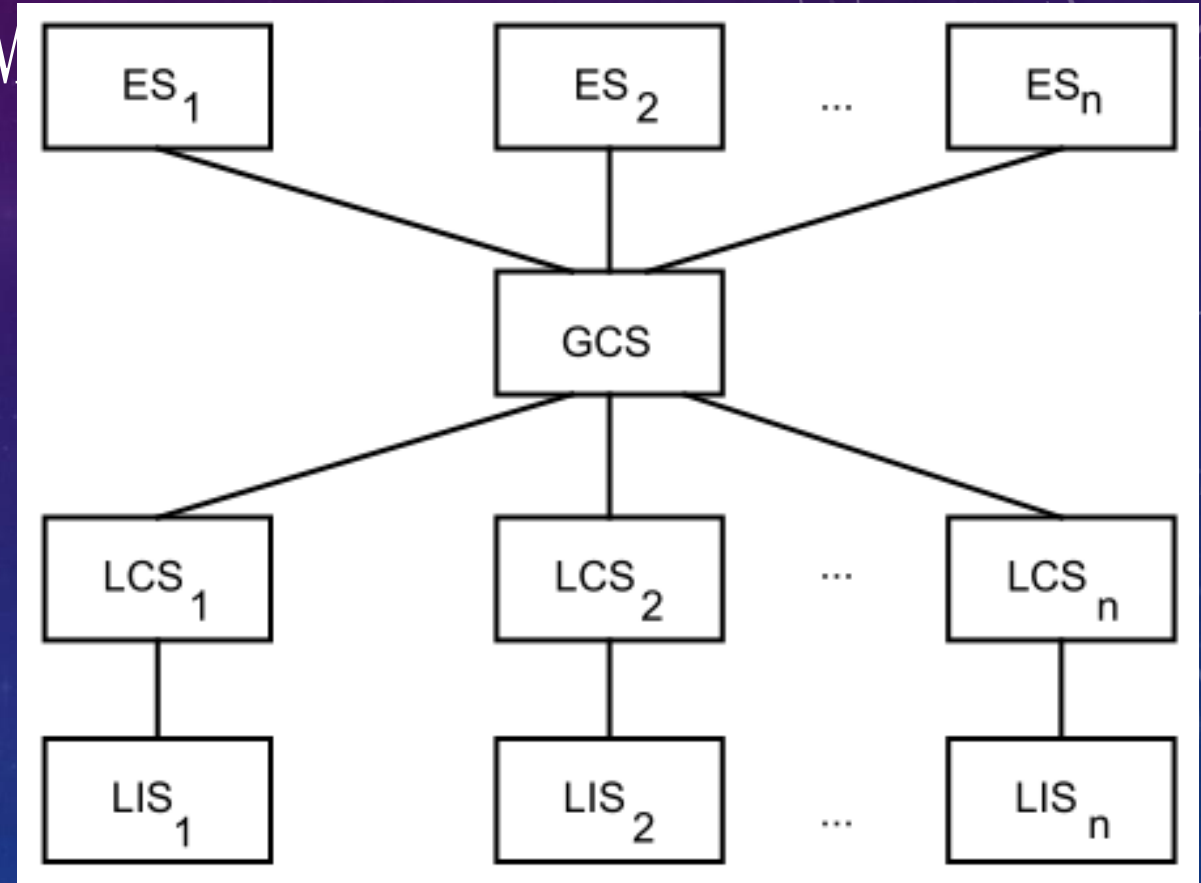
- Peer-to-peer (P2P) computing or networking is a distributed application architecture that partitions tasks or workloads between peers.
- Peers are **equally** privileged, equipotent participants in the application. They are said to form a peer-to-peer network of nodes.
- Peers make a portion of their resources, such as processing power, disk storage or network bandwidth, directly available to other network participants, without the need for central coordination by servers or stable hosts.



# DISTRIBUTED DBMS ARCHITECTURE

## – PEER-TO-PEER SYSTEM

- Local internal schema, LIS  
内部模式
- Global conceptual schema, GCS  
全局概念模式
- External schema, ES  
外部模式



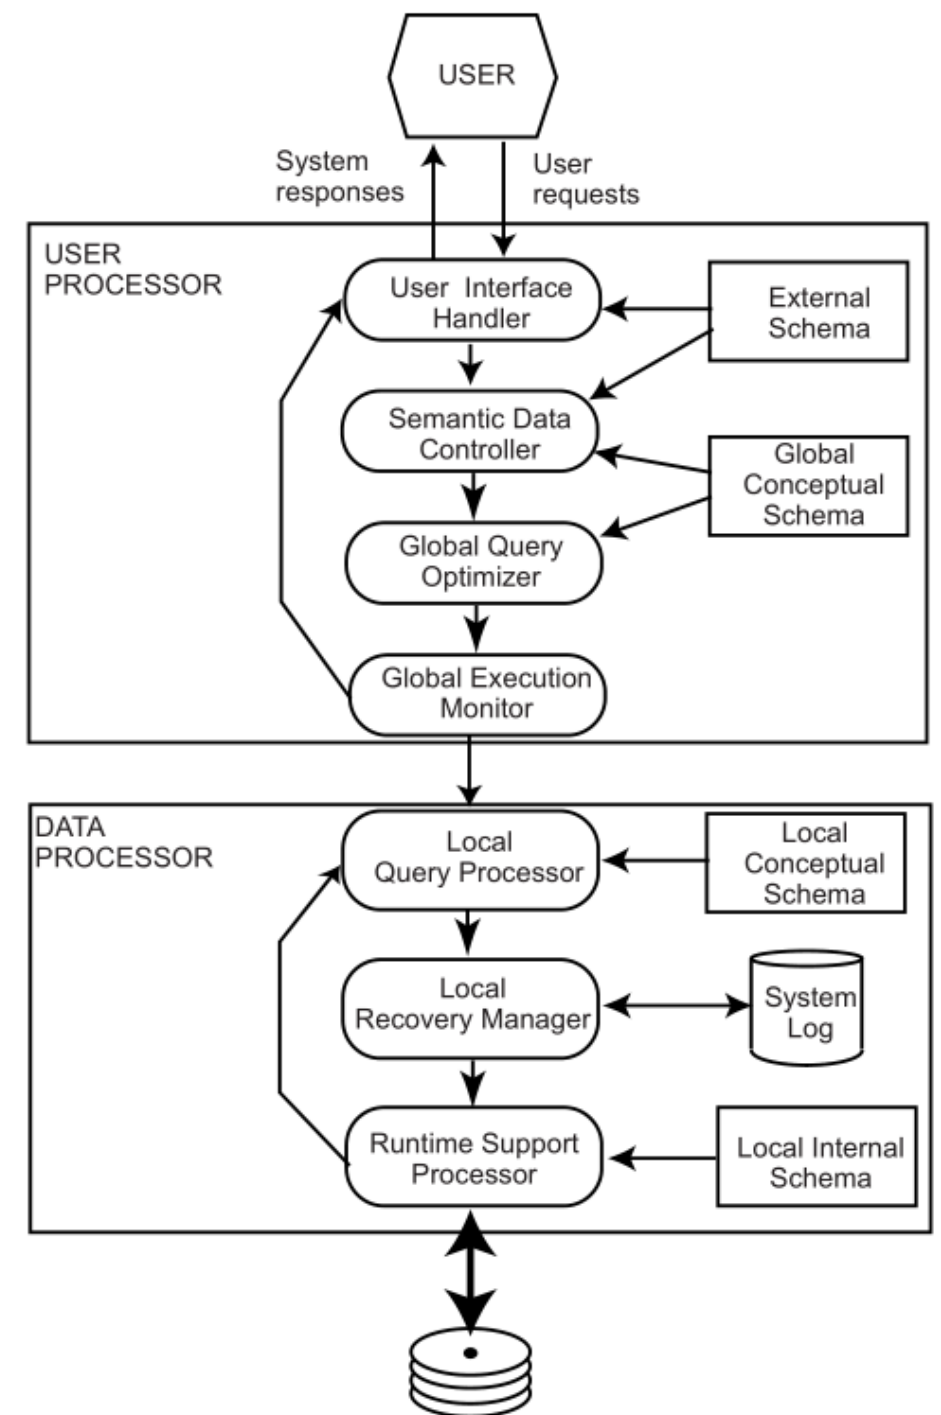
Distributed Database Reference Architecture



# DISTRIBUTED DBMS ARCHITECTURE

## – PEER-TO-PEER SYSTEMS

- (1) the user processor
- (2) the data processor



# DISTRIBUTED DBMS ARCHITECTURE

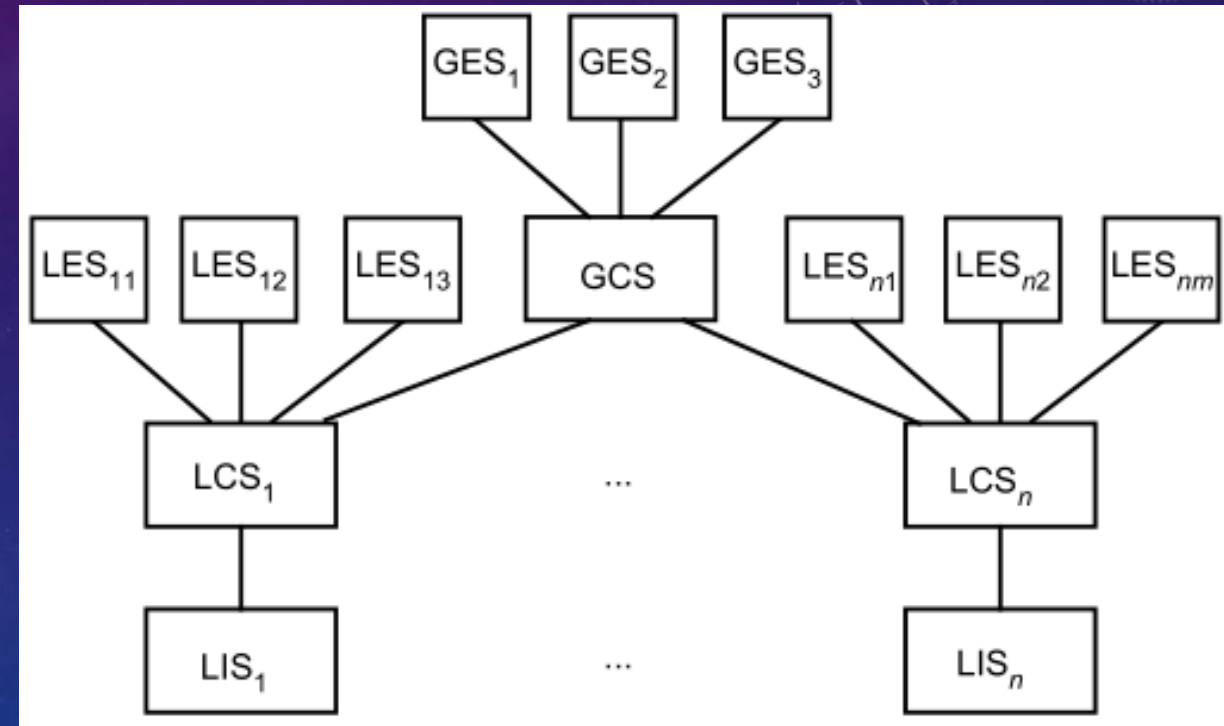
## – MULTIDATABASE SYSTEM ARCHITECTURE

- Multi-database systems (MDBS)
  - It represent the case where individual DBMSs (whether distributed or not) are fully autonomous and have no concept of cooperation;
  - they may not even “know” of each other’s existence or how to talk to each other.
- Export schema, ES出口模式

# DISTRIBUTED DBMS ARCHITECTURE

## – MULTIDATABASE SYSTEM ARCHITECTURE

- (1) the distributed multi-DBMSs
- (2) distributed DBMSs
- The differences in the level of autonomy between the distributed multi-DBMSs and distributed DBMSs are also reflected in their architectural models



MDBS Architecture with a GCS  
Global Conceptual Schema




# DISTRIBUTED DBMS ARCHITECTURE

## – MULTIDATABASE SYSTEM ARCHITECTURE

- If heterogeneity exists in the system, then two implementation alternatives exist: **unilingual**单语言 and **multilingual**多语言
  - A unilingual multi-DBMS requires the users to utilize possibly different data models and languages when both a local database and the global database are accessed.
  - the multilingual architecture permits each user to access the global database (i.e., data from other databases) by means of an external schema, defined using the language of the user's local DBMS.

# DISTRIBUTED DBMS ARCHITECTURE

- MULTIDATABASE SYSTEM ARCHITECTURE



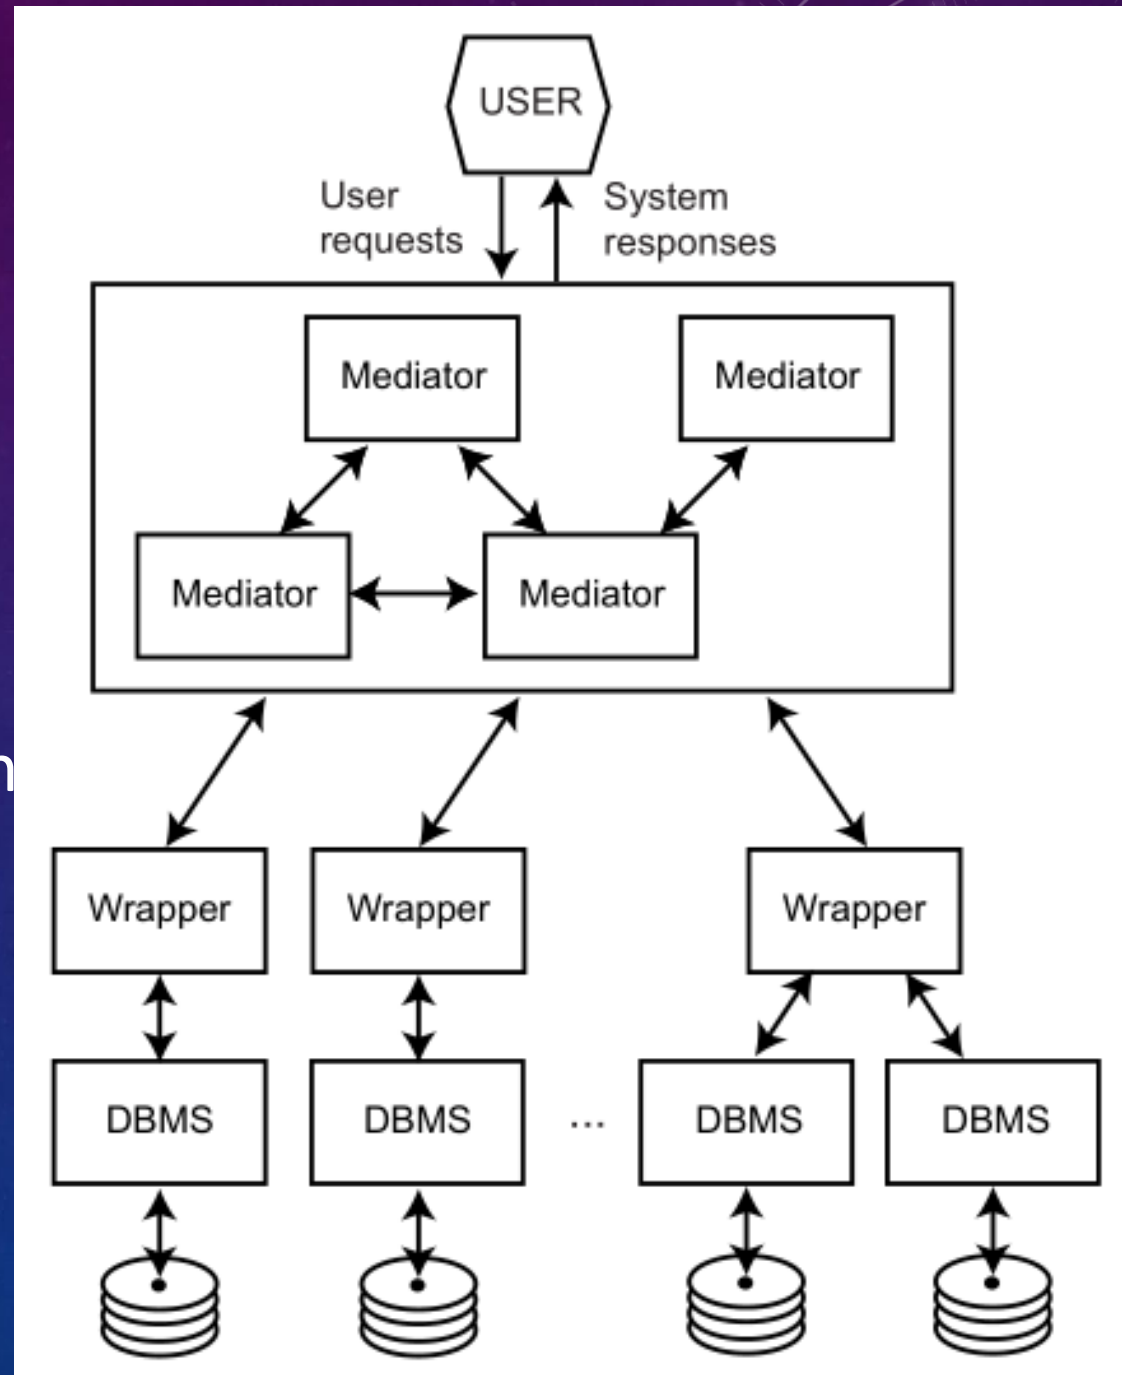
Components of an MDBS

# DISTRIBUTED DBMS ARCHITECTURE

## – MULTIDATABASE SYSTEM


## ARCHITECTURE

- Mediator中介程序
  - a software module that exploits encoded knowledge about certain sets or subsets of data to create information for a higher layer of applications
- Wrapper包装程序
  - whose task is to provide a mapping between a source DBMSs view and the mediators' view






# NETWORK



# NETWORK



# TYPES OF NETWORKS

- Scale
  - LAN, WAN, MAN
- Topology
  - star, ring, complete(mesh)
- Communication Schemes
  - point-to-point, broadcast (multi-point)
  - Microwave微波, cellular蜂窝, wireless broadband



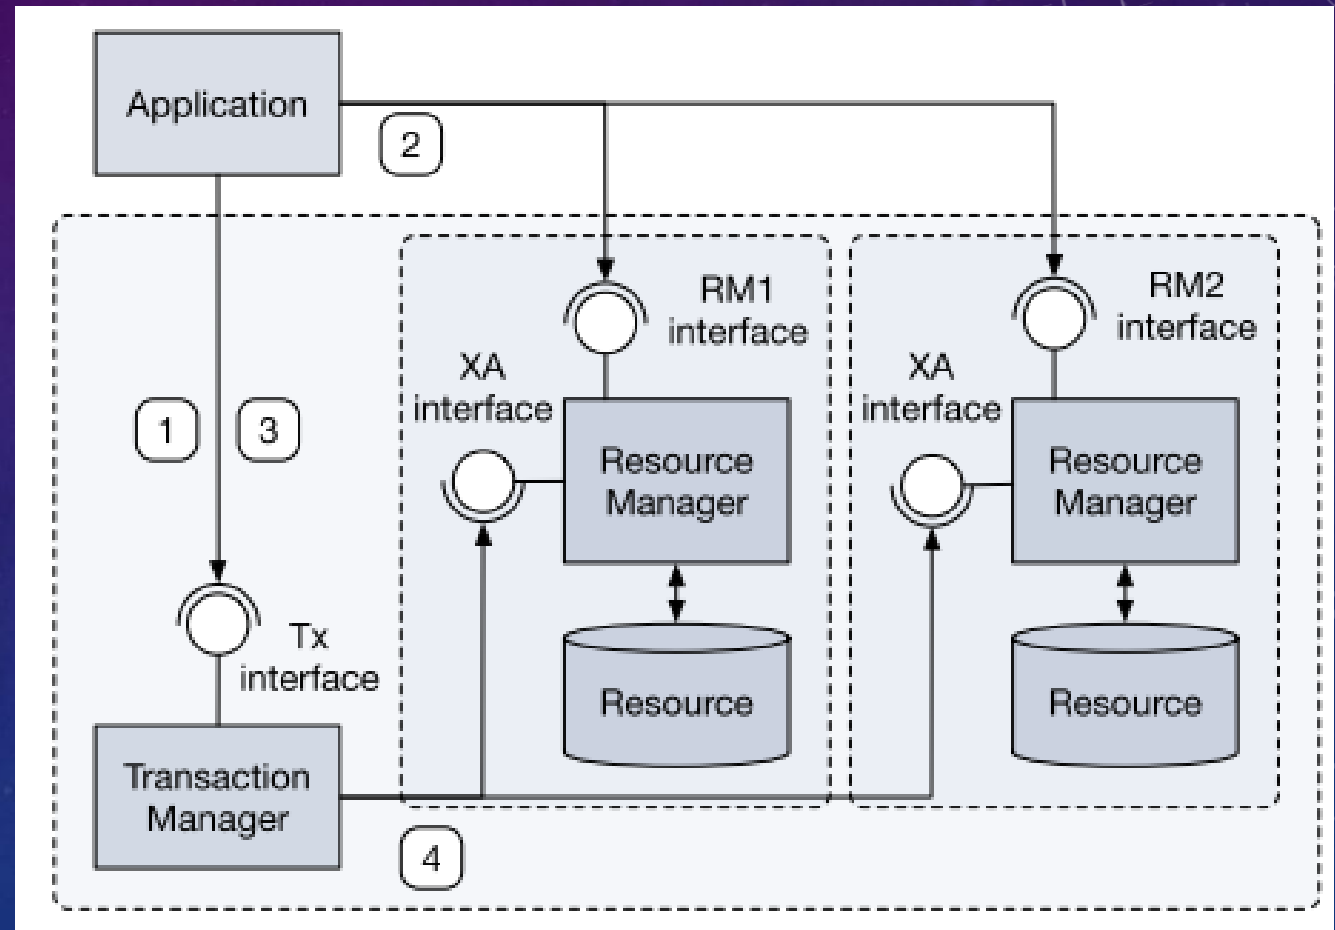
# THEOREMS OF DISTRIBUTED DB

- ACID theorem
- CAP theorem
- BASE theorem

# ACID THEOREM

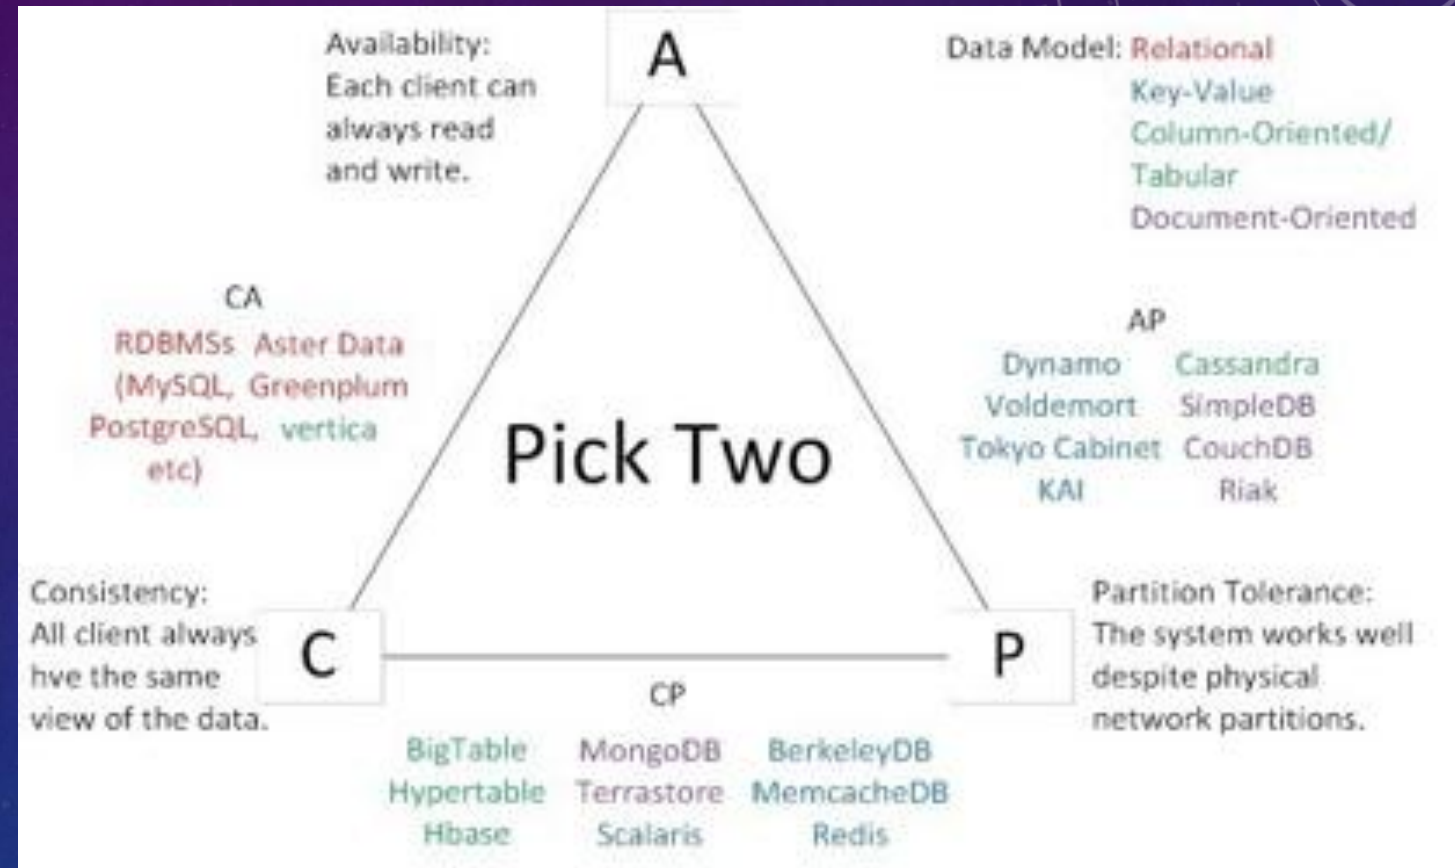
- Atomicity
- Consistency
- Isolation
- Durability

## Transaction Processing System



# CAP THEOREM

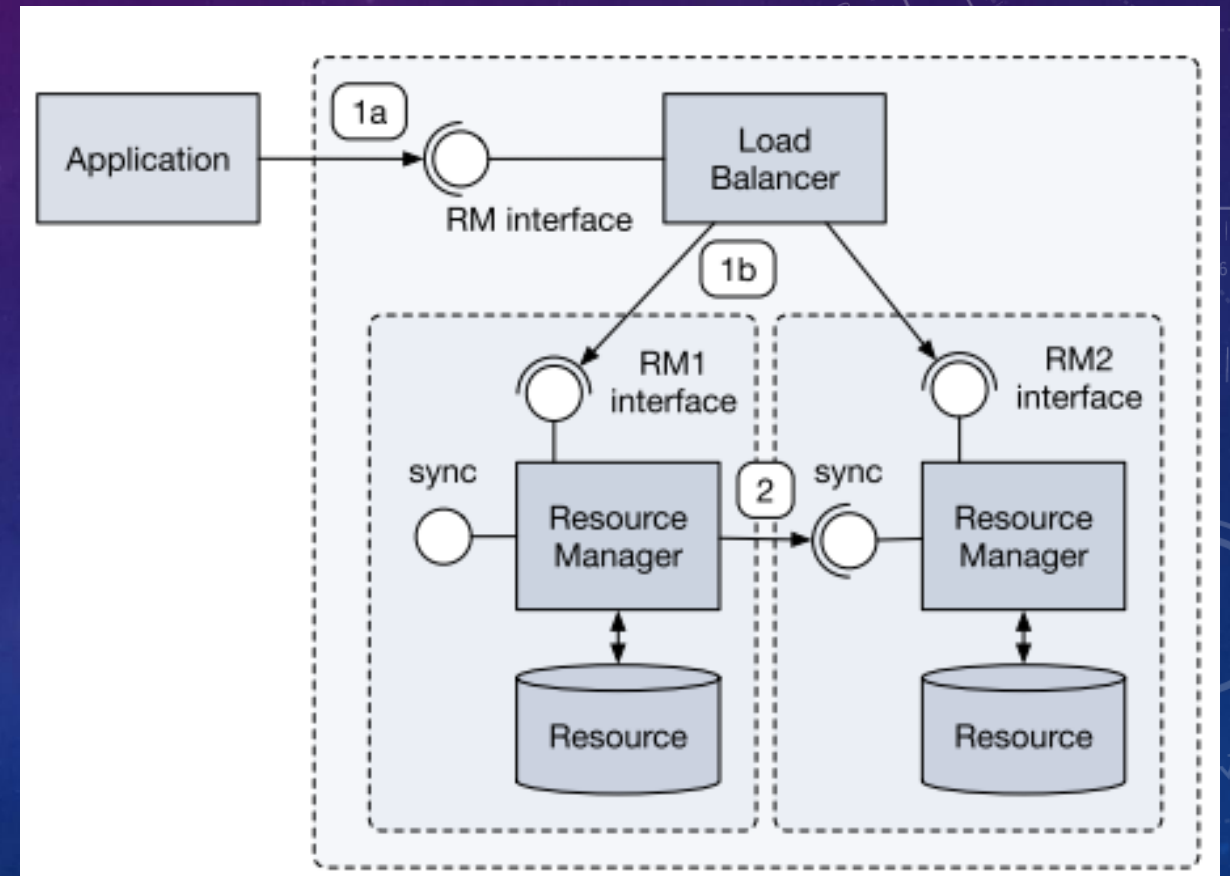
it is **impossible** for a distributed data store to simultaneously provide more than two out of the following three guarantees



# BASE THEOREM


- Basically Available
- Soft state
- **Eventually** consistent

## App Based BASE MODEL





# Data as a Service Data as a Platform



Unit of Storage	Desc
Byte 字节	1 Byte = 8 bit
KB (Kilobyte, 千字节)	1KB=1024 Byte
MB (Megabyte, 兆字节)	1MB=1024KB
GB (Gigabyte, 吉字节)	1GB=1024MB
TB (Trillionbyte, 太字节)	1TB=1024GB
PB (Petabyte, 拍字节)	1PB=1024TB
EB (Exabyte, 艾字节)	1EB=1024PB
ZB (Zettabyte, 泽字节)	1ZB=1024EB
YB(Yottabyte, 尧字节)	1YB=1024ZB

# DATA MINING

- Data Mining

- Refers loosely to the process of semi-automatically analyzing large DB to find useful patterns
  - Case: beer and nappy
- Knowledge discovery in artificial intelligence: machine learning



# DATA WAREHOUSE

- Data warehouse
  - 数据仓库
  - Gather data from multiple sources under a unified schema, at a single site
  - Then, they provide the user a single uniform interface to data



# INFORMATION RETRIEVAL

- Information Retrieval 信息检索
  - Querying of unstructured textual data
  - Textual data is unstructured, unlike the rigidly structured data in relational databases
  - Textual data has grown explosively

# 2016E = Milestone Year for 'Traditional' Live Streaming on Social Networks... NFL Live Broadcast TV of Thursday Night Football on Twitter (Fall 2016)

## INFORMATION RETRIEVAL: TODAY

### *Hypothetical Mock-Up*

Complete Sports Viewing Platform =

Live Broadcast + Analysis + Scores + Replays + Notifications + Social Media Tools

*Tune-In Notifications,  
Game Reminders,  
Breaking Actions*



*Scoreboard Allows Fans to  
Follow Play-by-Play*



*Tweets Engage Fans in  
Real-Time Conversation*



*Professional  
Commentary and  
Analysis*



*Vertical View =  
Live Broadcast + Tweets  
Dashboard for Social  
Media Engagement*



*Toggle Between Tweets  
from Stadium / Nearby / All*

*Horizontal View =  
Unencumbered, Full-  
Screen, TV-like Viewing  
Experience*



Next ...



# RANK OF DATABASES

354 systems in ranking, March 2020

Rank			DBMS	Database Model	Score		
Mar 2020	Feb 2020	Mar 2019			Mar 2020	Feb 2020	Mar 2019
1.	1.	1.	Oracle +	Relational, Multi-model ⓘ	1340.64	-4.11	+61.50
2.	2.	2.	MySQL +	Relational, Multi-model ⓘ	1259.73	-7.92	+61.48
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model ⓘ	1097.86	+4.11	+50.01
4.	4.	4.	PostgreSQL +	Relational, Multi-model ⓘ	513.92	+6.98	+44.11
5.	5.	5.	MongoDB +	Document, Multi-model ⓘ	437.61	+4.28	+36.27
6.	6.	6.	IBM Db2 +	Relational, Multi-model ⓘ	162.56	-2.99	-14.64
7.	7.	↑ 9.	Elasticsearch +	Search engine, Multi-model ⓘ	149.17	-2.98	+6.38
8.	8.	8.	Redis +	Key-value, Multi-model ⓘ	147.58	-3.84	+1.46
9.	9.	↓ 7.	Microsoft Access	Relational	125.14	-2.92	-21.07
10.	10.	10.	SQLite +	Relational	121.95	-1.41	-2.92

- Relational DBMS
- Key-value stores
- Document stores
- Graph DBMS
- Time Series DBMS
- Object oriented DBMS
- RDF stores
- Search engines
- Wide column stores
- Multivalue DBMS
- Native XML DBMS
- Event Stores
- Content stores
- Navigational DBMS



# RESEARCH TOPICS

数据库实现新技术	云计算环境中的数据管理
Web数据管理	查询处理与查询优化
数据流管理	XML和半结构化数据
数据仓库和OLAP	近似和非确定性数据库
内容与知识管理	数据挖掘和知识发现
元数据管理	数据集成和迁移
嵌入式数据库与移动数据库	并行和分布式数据库系统
特定领域的数据库系统	数据库自我管理
智能用户接口技术	空间和时态数据库系统
多媒体数据库技术	数据隐私与安全
信息检索与数据库	协同工作技术
物联网数据管理	

# OBJECTIVES

- The Age of Database
- Distributed Data Processing
- What is a Distributed Database System?
- Data Delivery Alternatives
- Promises of DDBSs
- Complications Introduced by Distribution
- Design Issues
- Distributed DBMS Architecture

# HISTORY OF DATABASE SYSTEMS

- 1950s and early 1960s: tape
- Late 1960s and 1970s
  - Hard disk, file, DB(hierarchical, network)
  - network db
    - CODASYL ,Integrated Data Store (IDS)
- **Codd**, E. F.. "A relational model of data for large shared data banks." Communications of The ACM 13.6 (1970): 377-387.




# HISTORY OF DATABASE SYSTEMS

- 1980s
  - System R: IBM
    - Astrahan, Morton M., et al. "System R: relational approach to database management." ACM Transactions on Database Systems 1.2 (1976): 97-137.
  - Ingres: BSD
  - IBM DB2, Oracle, DEC Rdb
- Early 1990s
  - Object-relational DB



# EDGAR F. CODD

- Edgar Frank Codd
  - Edgar Frank "Ted" Codd was an English computer scientist who, while working for IBM, invented the relational model for database management, the theoretical basis for relational databases
  - ACM Turing Award, 1981
  - Online analytical processing (OLAP)
  - 19 August 1923 – 18 April 2003



# HISTORY OF DATABASE SYSTEMS

- 1990s
  - WWW: DB had to support Web interface to data
- 2000s
  - XML, Xquery
  - PostgreSQL, MySQL
  - Giant data storage systems
    - Google BigTable, Yahoo PNuts, Amazon, ...

# SUMMARY

- The Age of Database
- Distributed Data Processing
- What is a Distributed Database System?
- Data Delivery Alternatives
- Promises of DDBSs
- Complications Introduced by Distribution
- Design Issues
- Distributed DBMS Architecture

• History

©LKD





Q&A?



# 赞数字文明时代之开启

李旭东 2017

文明初叶几时真，造化阴阳始幻尘。  
书简成山薪火旺，零壹遁迹智能春。  
有形百载多将朽，数字千年总是新。  
懵懂蹒跚别旧日，一朝奋起笑前津。





# THANKS !

leexudong@nankai.edu.cn